



System on Chip Design Lab.
晶片系統設計實驗室

Chapter 4 – part III

Advanced Pipelining Issues

Tien-Fu Chen

Dept. of Computer Science and
Information Engineering

National Chiao Tung Univ.

Material source:
RISC-V slides

Deal with Pipelining

❑ Advanced issues pipelining

- Instructions may have dependency
- Dependency may result into hazards
- More instruction-level parallelism:
 - Pipelining
 - Multiple issues
 - Multithreading

❑ What makes pipelining hard?

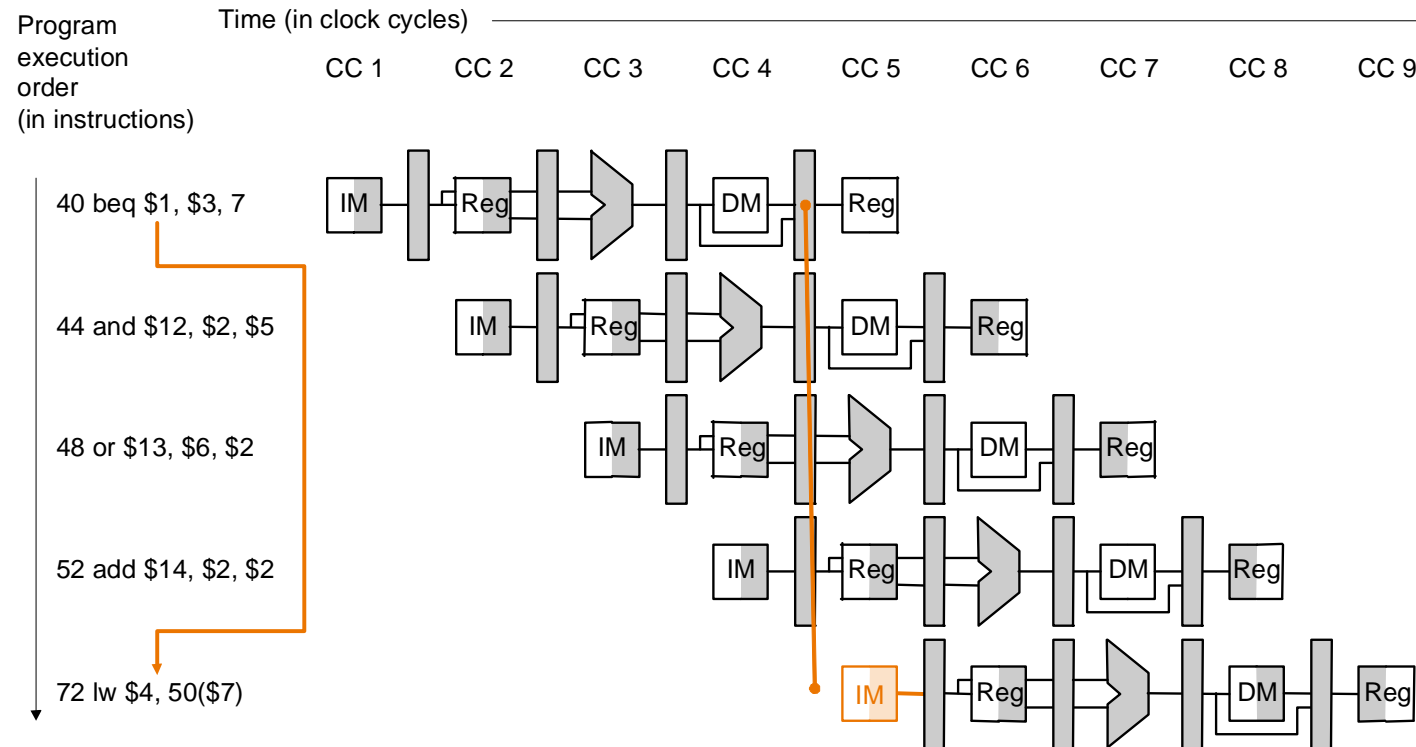
- structural hazards: resource conflicts, e.g. only one memory
 - sol: adding more resource
- control hazards: need to worry about branch instructions
 - sol: resolve earlier, predict branch
- data hazards: an instruction depends on a previous instruction
 - sol: forwarding, bypassing

Branch Hazards

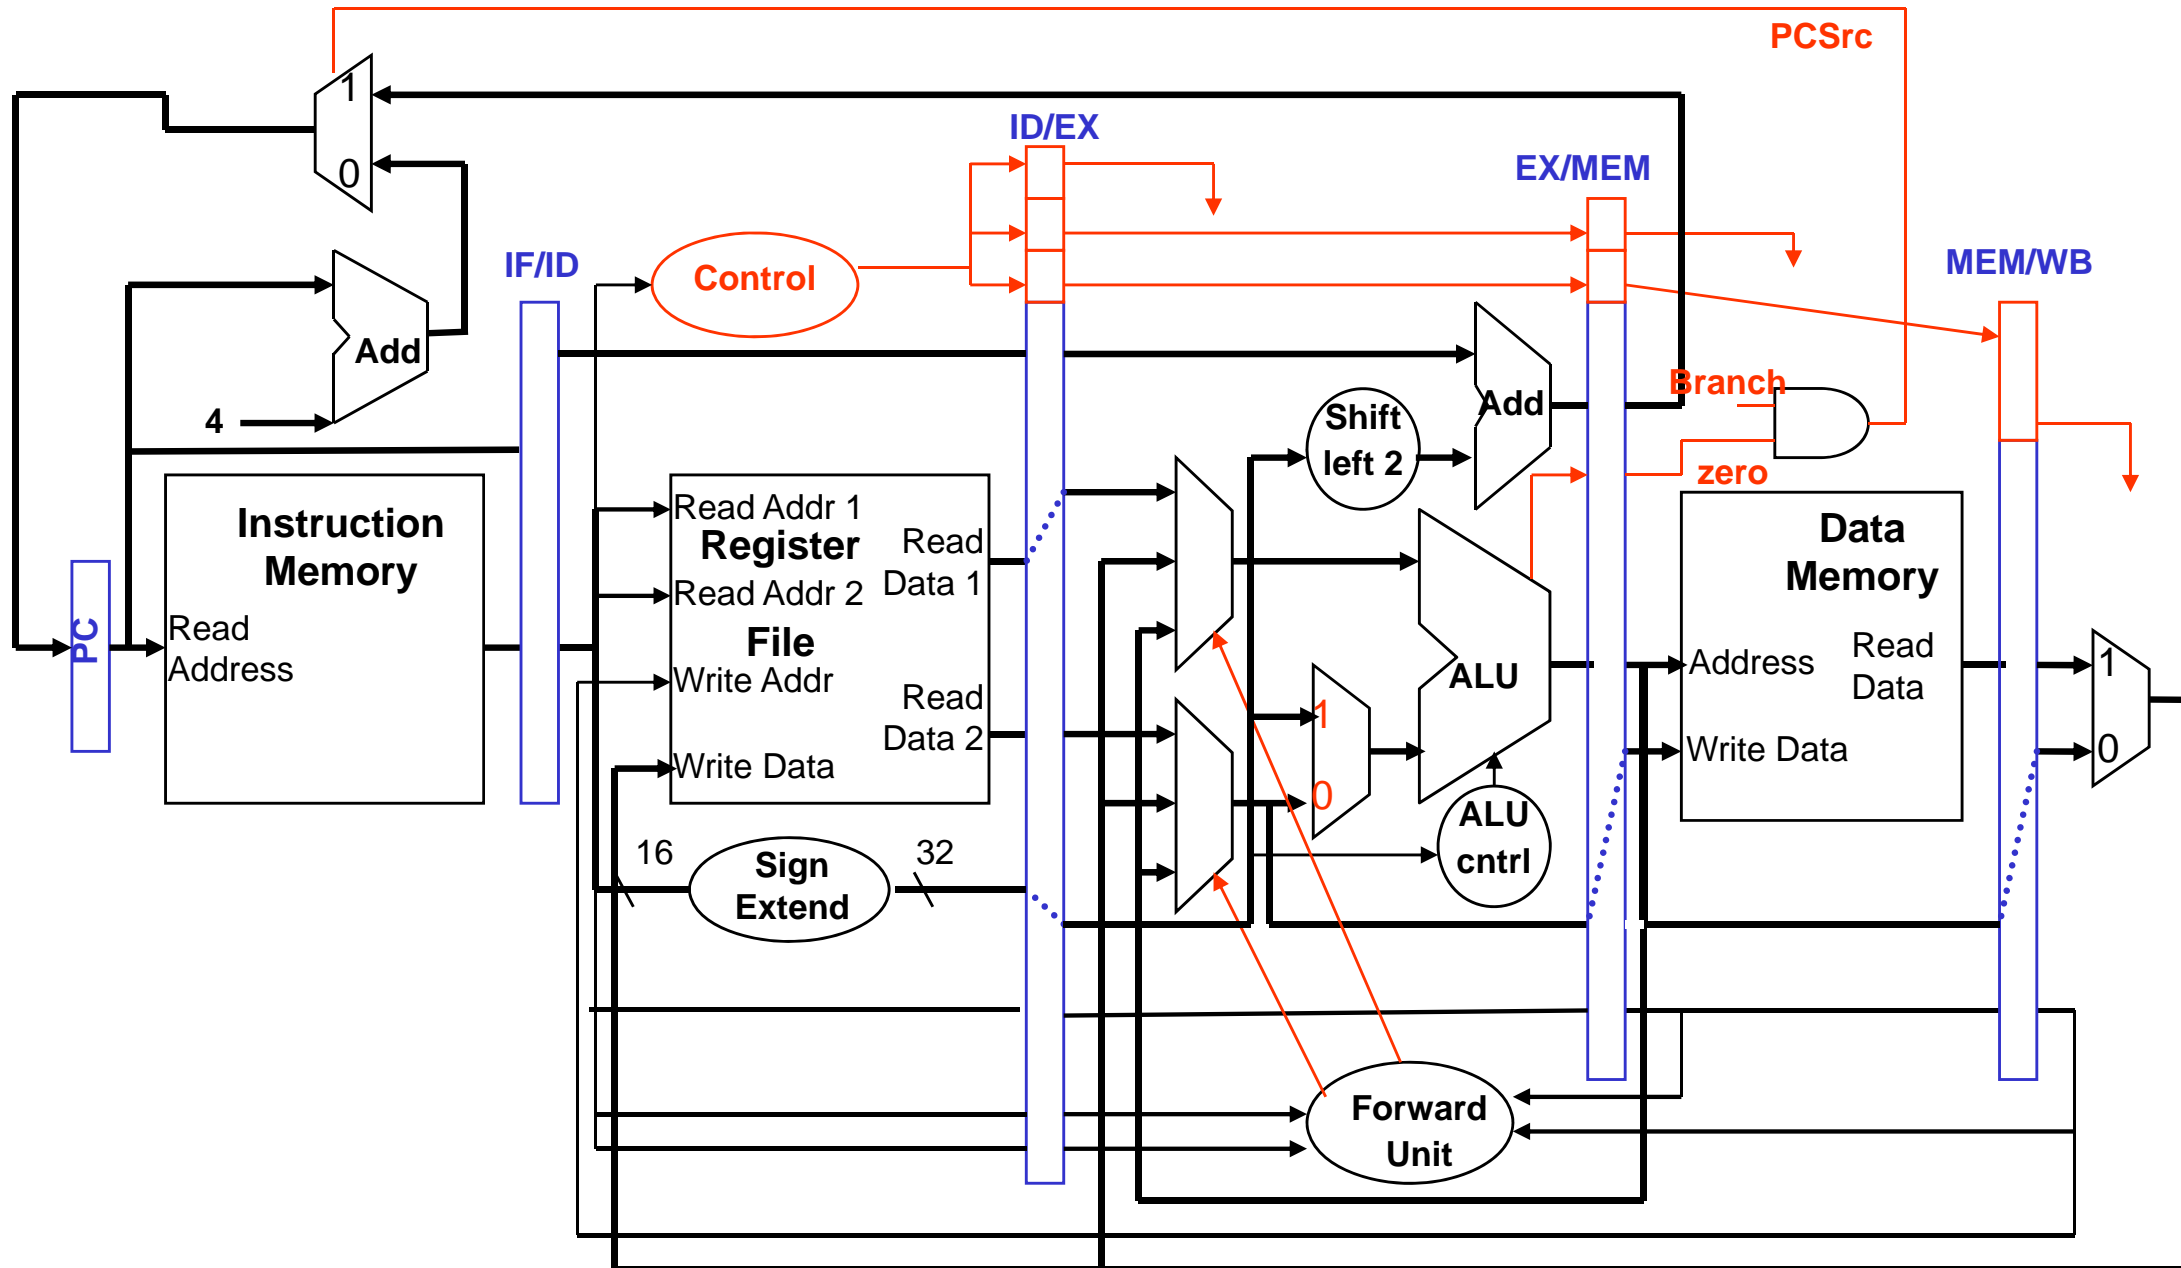
- ❑ 3 instructions after branch will be incorrectly executed

- ❑ Solutions:

- Stall
- Reducing branch delay by early determination of branch
- Instruction scheduling for branch hazard
- Branch condition prediction
 - 1-b buffer
 - 2-b buffer



3 cycles delay on a branch hazard



Moving Branch Decisions Earlier in Pipe

❑ Two key decisions:

- Target address adder
- Register comparator

❑ Move to the EX stage

- Reduces the number of stall cycles to two
- Adds an and gate and a 2x1 mux to the EX timing path

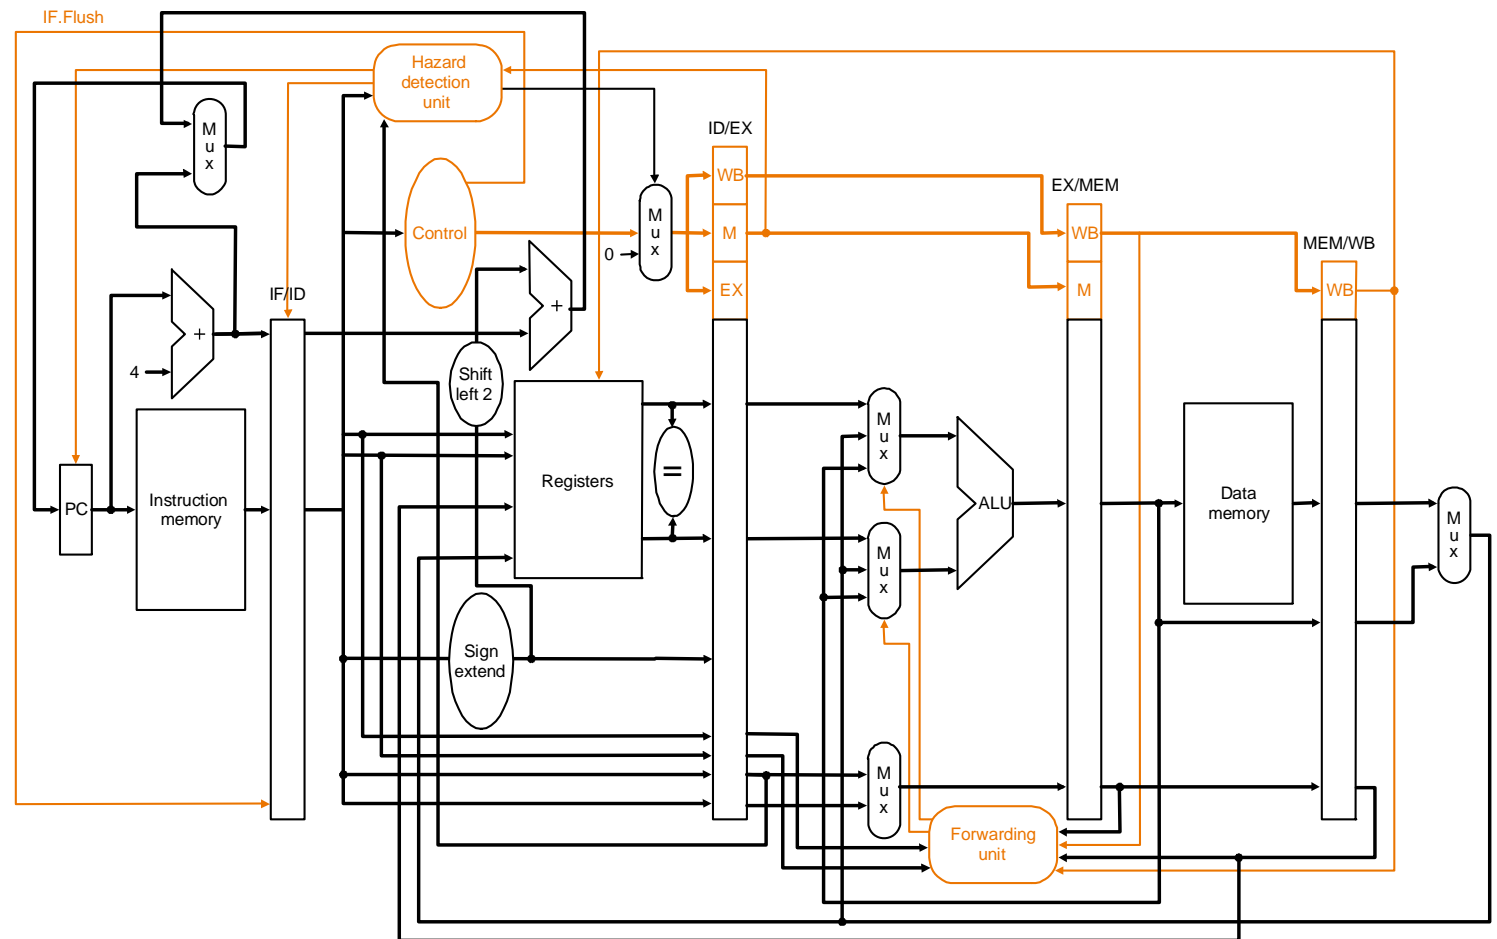
❑ Move to the ID stage

- Reduces the number of stall cycles to one (like with jumps)
- Computing branch target address can be done in parallel with RegFile read (done for all instructions – only used when needed)
- Comparing the registers can't be done until after RegFile read, so comparing and updating the PC adds a comparator, an and gate, and a 3x1 mux to the ID timing path
- Need forwarding hardware in ID stage

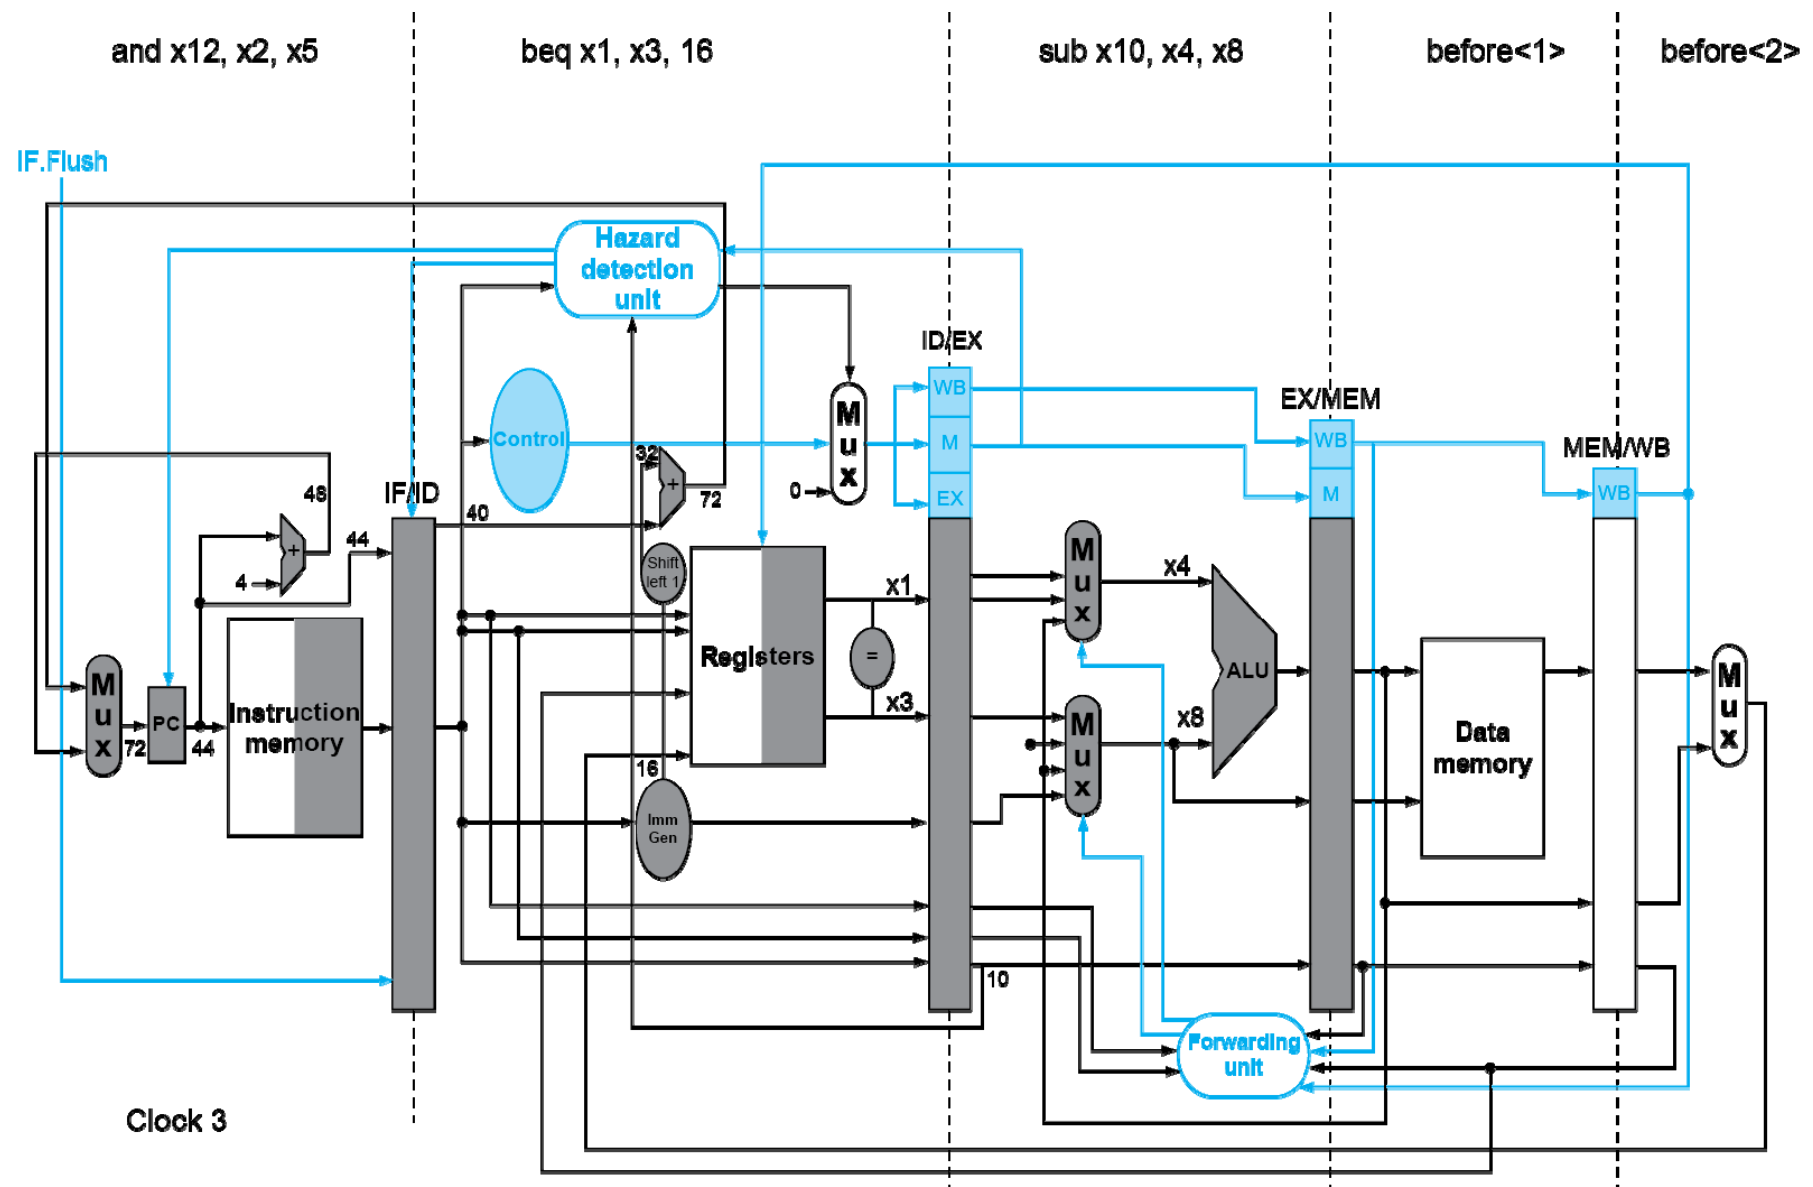
Flush instructions follows branch

❑ Move branch decision from Mem to ID

❑ Zeros the instruction for next cycle

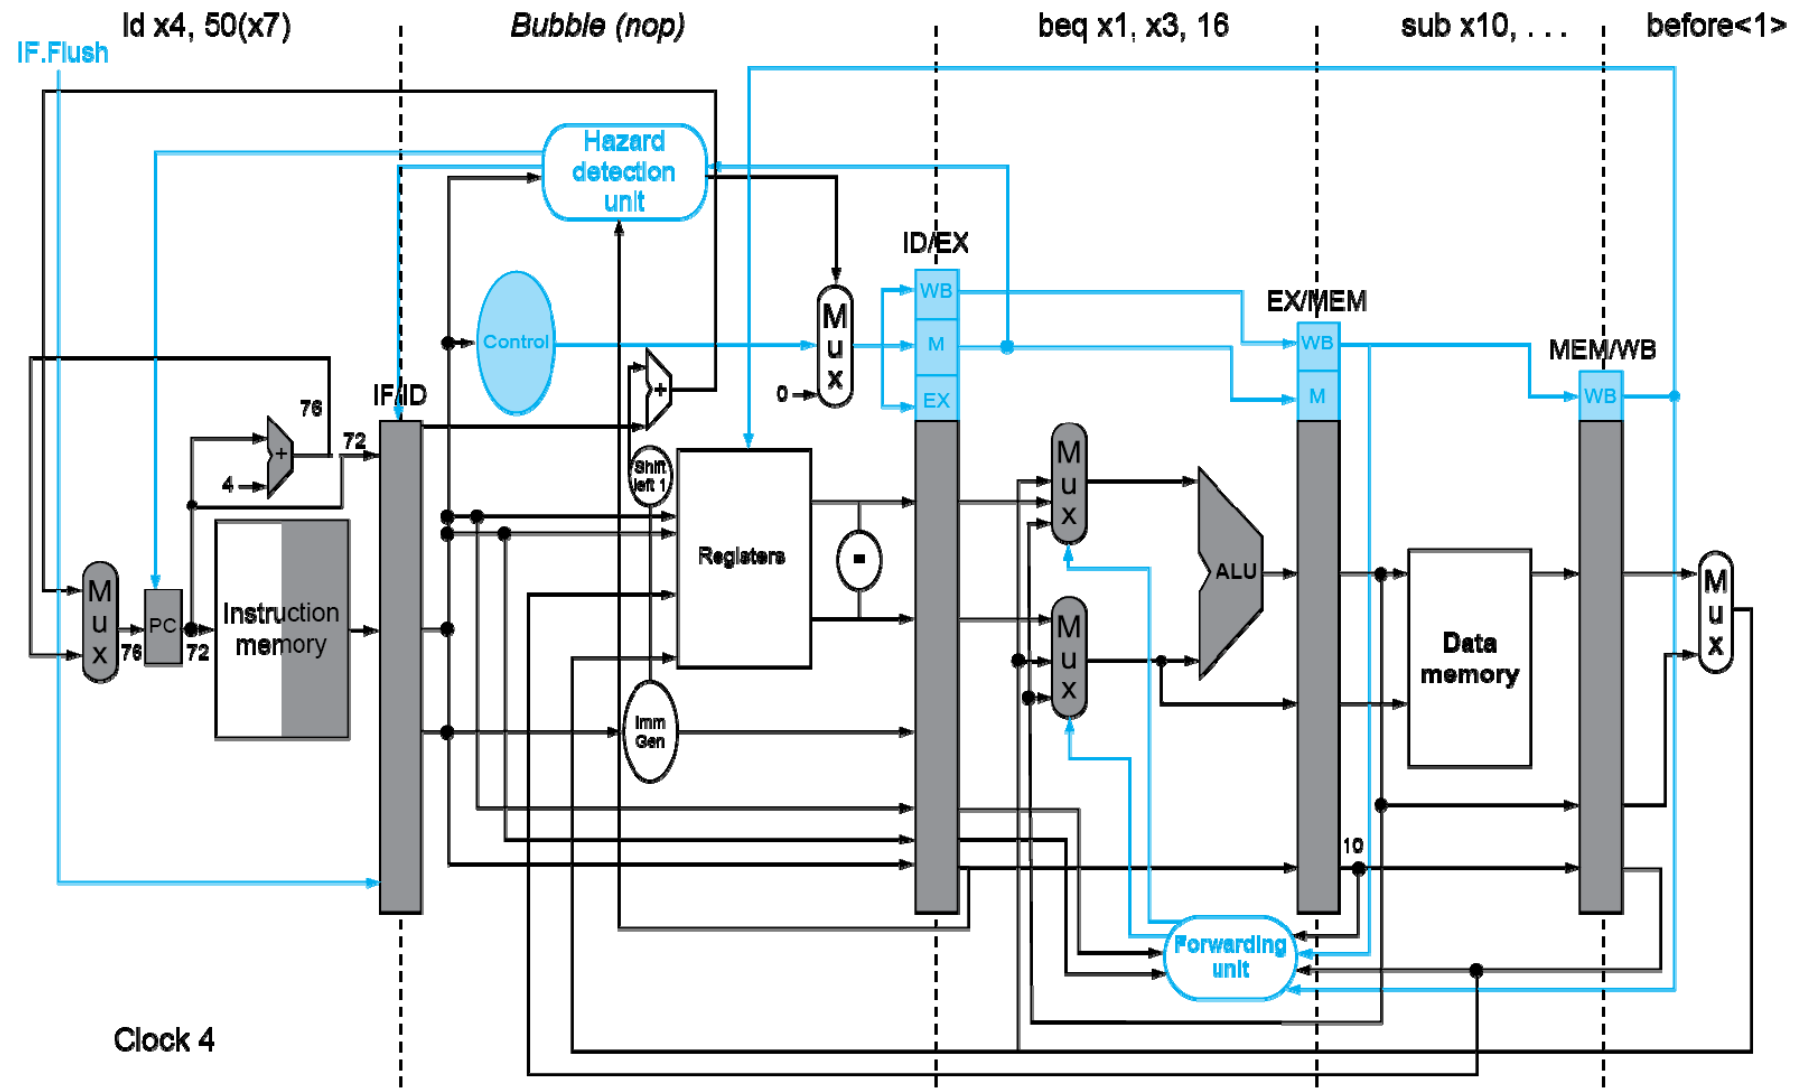


Example: Branch Taken

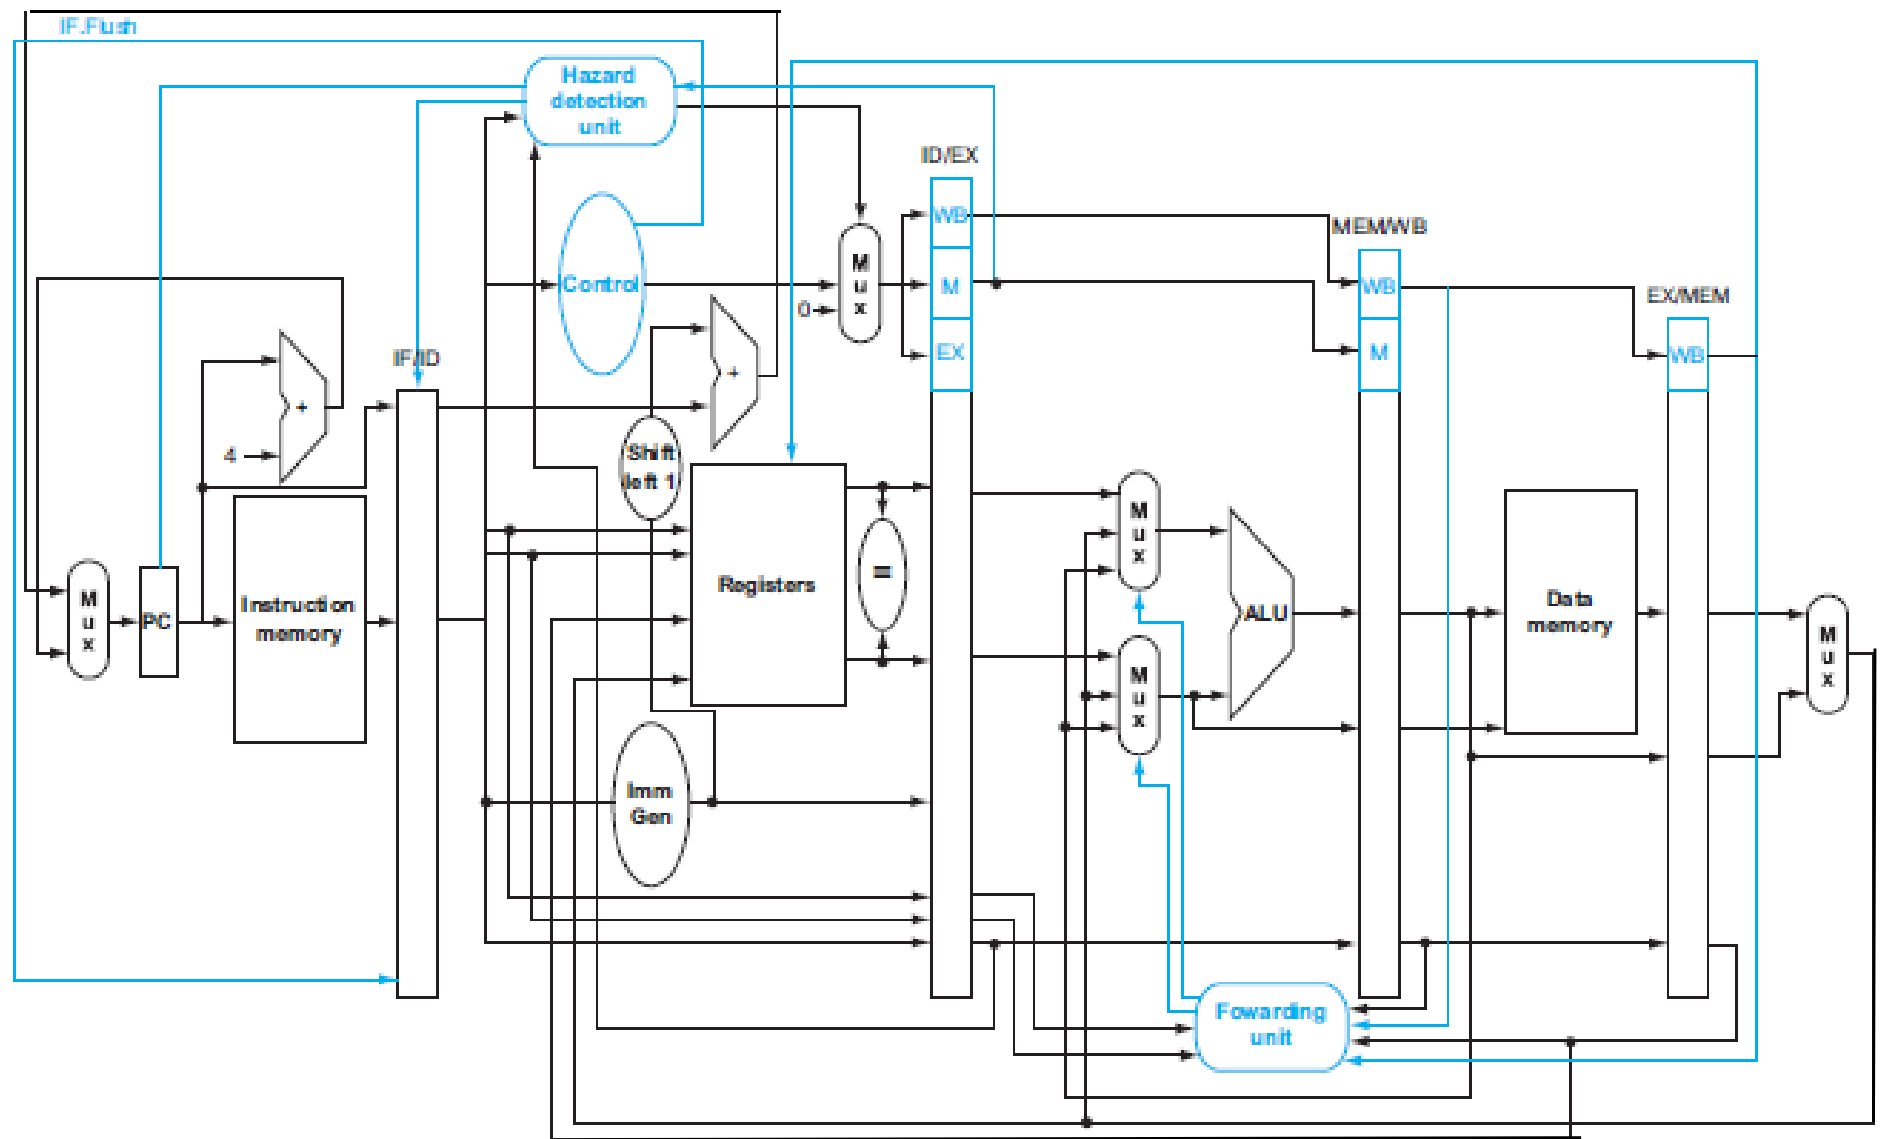


Clock 3

Example: Branch Taken



Datapath for branch

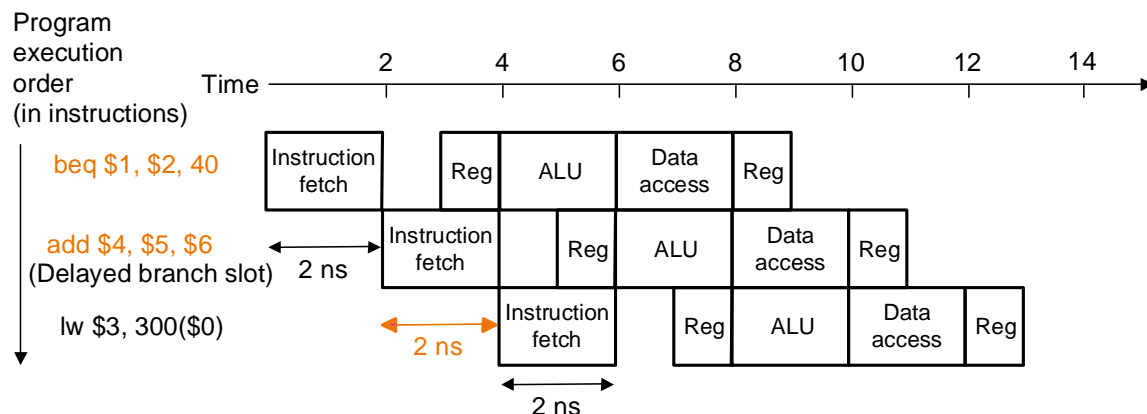


More solutions for control hazard

- ❑ Software solution by instruction scheduling
 - Three strategies to optimize branch

- ❑ Dynamic Branch Prediction
 - dynamic branch prediction
 - predict taken/not taken
 - Predict branch target

- ❑ Delayed branch

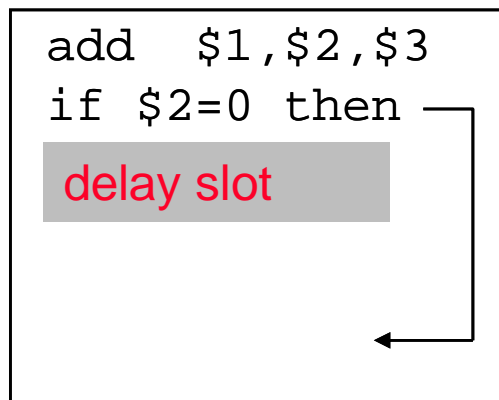


- ❑ Conditional instruction set
 - 4-b condition in 32b ARM instructions


Software solutions for branch hazards

- ❑ Software by instruction scheduling
 - branch delay slot
 - moves an instruction to immediately after the branch that is not affected by the branch (a **safe** instruction) thereby **hiding** the branch delay

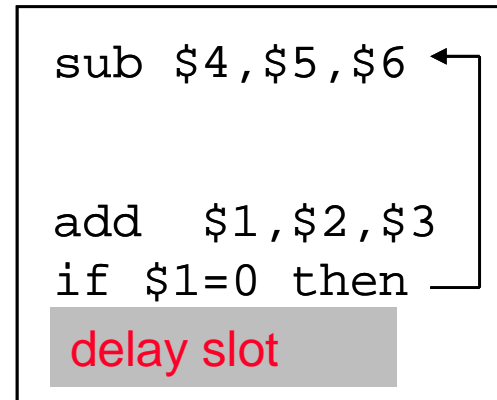
A. From before branch




becomes



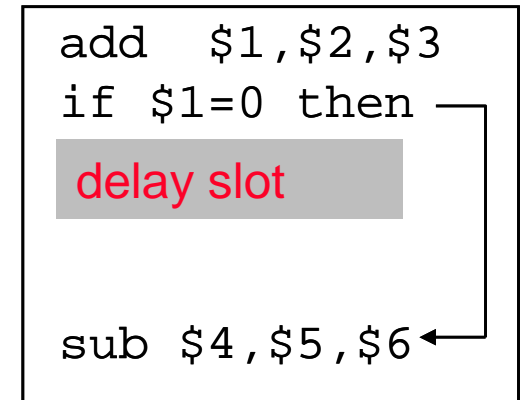
B. From branch target



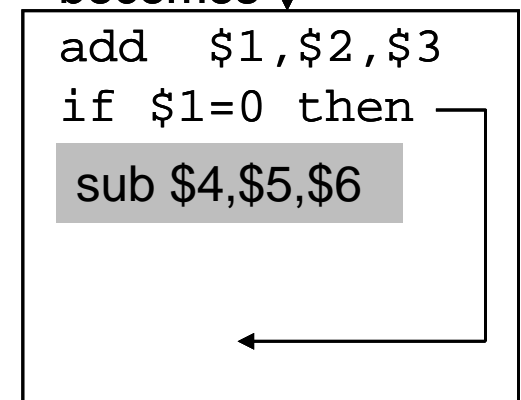
becomes



C. From fall through



becomes



Branch Condition Prediction

- ❑ Predict branch direction: taken or not taken (T/NT)



- ❑ Static prediction: compilers decide the direction
- ❑ Dynamic prediction: hardware decides the direction using dynamic information
 1. 1-bit Branch-Prediction Buffer
 2. 2-bit Branch-Prediction Buffer
 3. Correlating Branch Prediction Buffer
 4. Tournament Branch Predictor
 5. and more ...
- ❑ Branch target prediction: branch target buffer (BTB)

Dynamic Branch Prediction


- ❑ Deeper and superscalar pipelines produce larger branch penalty
- ❑ A **branch prediction buffer** (aka branch history table (**BHT**))

In the IF stage addressed by the lower bits of the PC, contains bit(s) passed to the ID stage through the IF/ID pipeline register that tells whether the branch was taken last time it was executed


- ❑ Use dynamic prediction
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

Predictor for a Single Branch

General Concept




1-bit prediction



Branch History Table **BHT** of 1-bit Predictor

- ❑ Can use only one 1-bit predictor, but accuracy is low
- ❑ BHT: use a table of simple predictors, indexed by bits from PC
- ❑ Similar to direct mapped cache
- ❑ More entries, more cost, but less conflicts, higher accuracy
- ❑ BHT can contain complex predictors




1-bit BHT Weakness

- ❑ Example: in a loop, 1-bit BHT will cause 2 mispredictions
- ❑ Consider a loop of 9 iterations before exit:

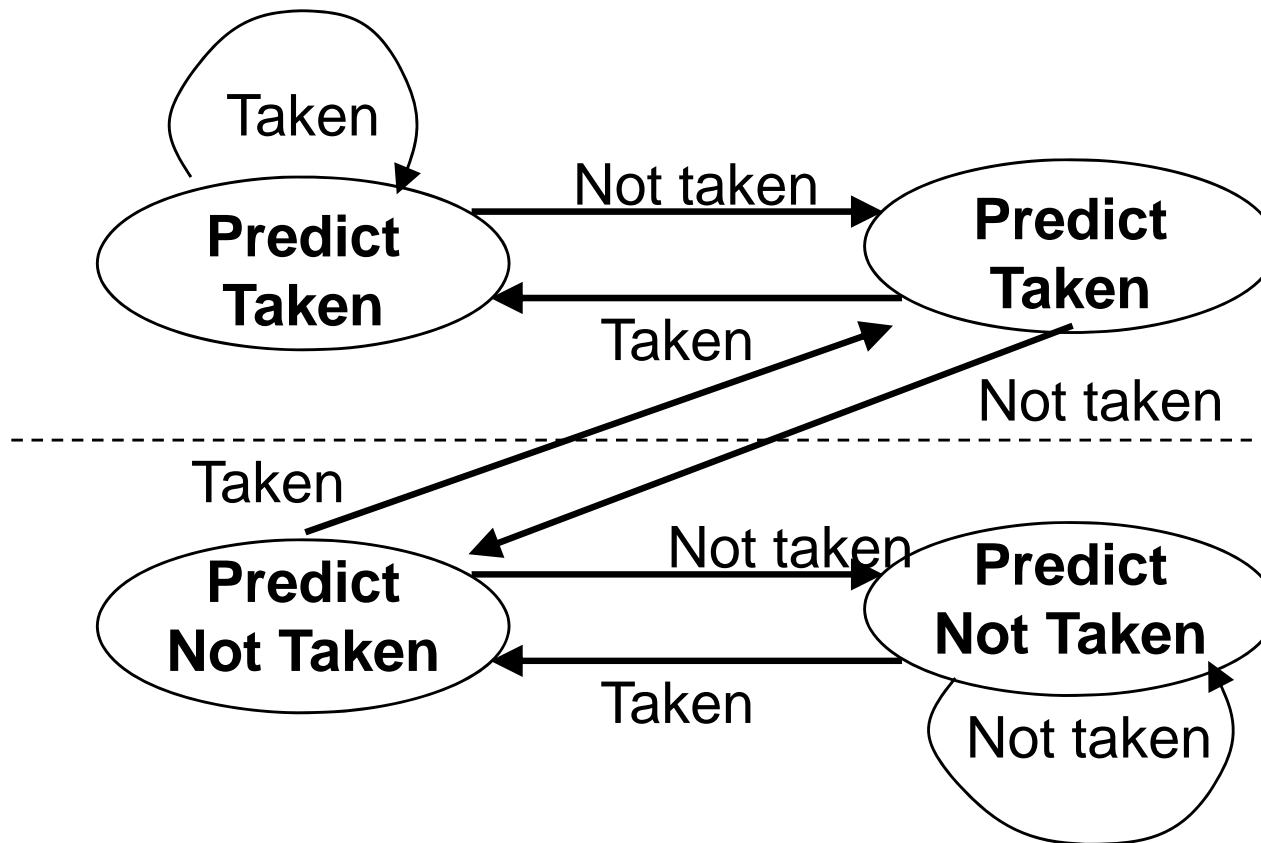
```
for (...) {  
    for (i=0; i<9; i++)  
        a[i] = a[i] * 2.0;  
}
```

- End of loop case, when it exits instead of looping as before
- First time through loop on *next* time through code, when it predicts *exit* instead of looping
- Only 80% accuracy even if loop 90% of the time



2-bit Predictors

- ❑ A 2-bit scheme can give 90% accuracy since a prediction must be wrong twice before the prediction bit is changed

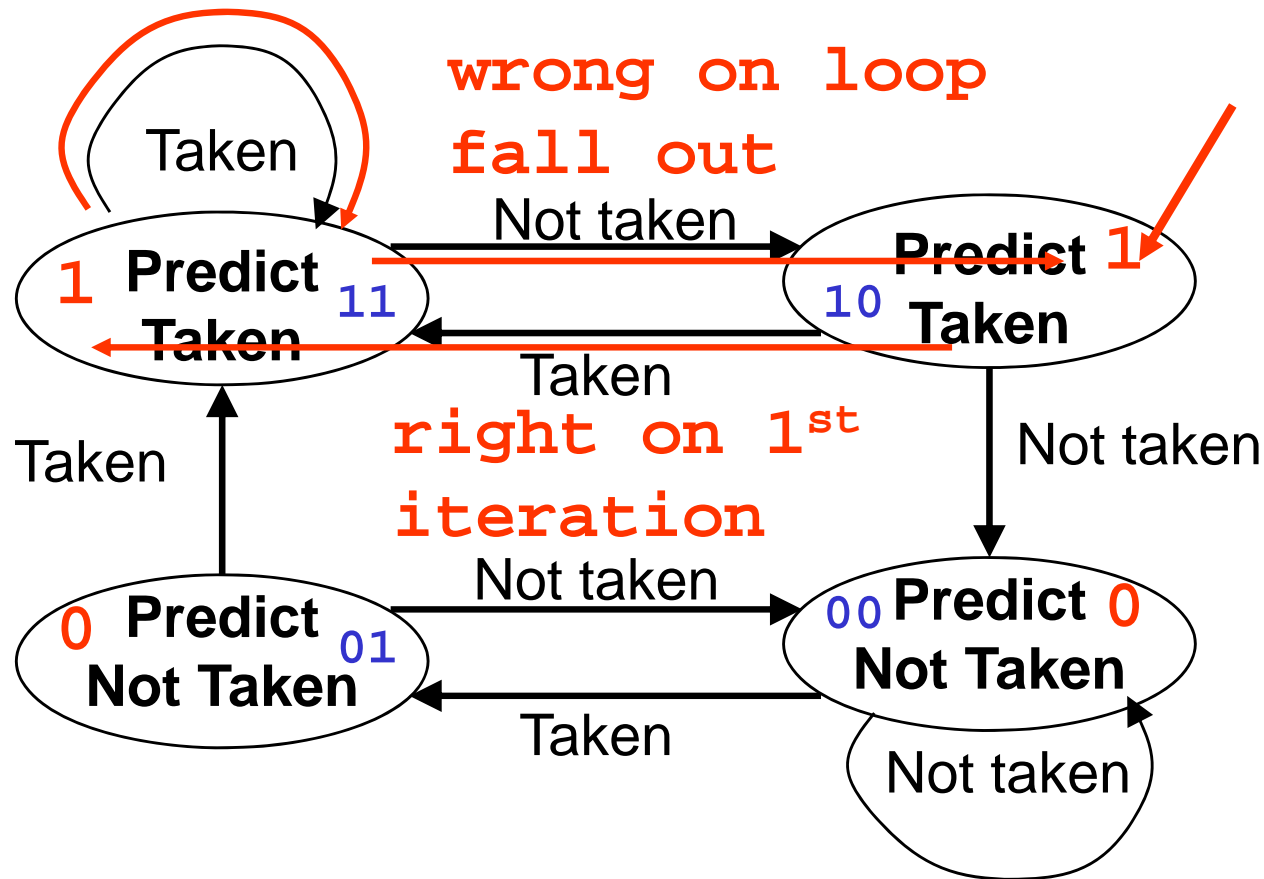


```
Loop: 1st loop instr
      2nd loop instr
      .
      .
      .
      last loop instr
      bne $1,$2,Loop
      fall out instr
```

2-bit Predictors

- A 2-bit scheme can give 90% accuracy since a prediction must be wrong twice before the prediction bit is changed

right 9 times



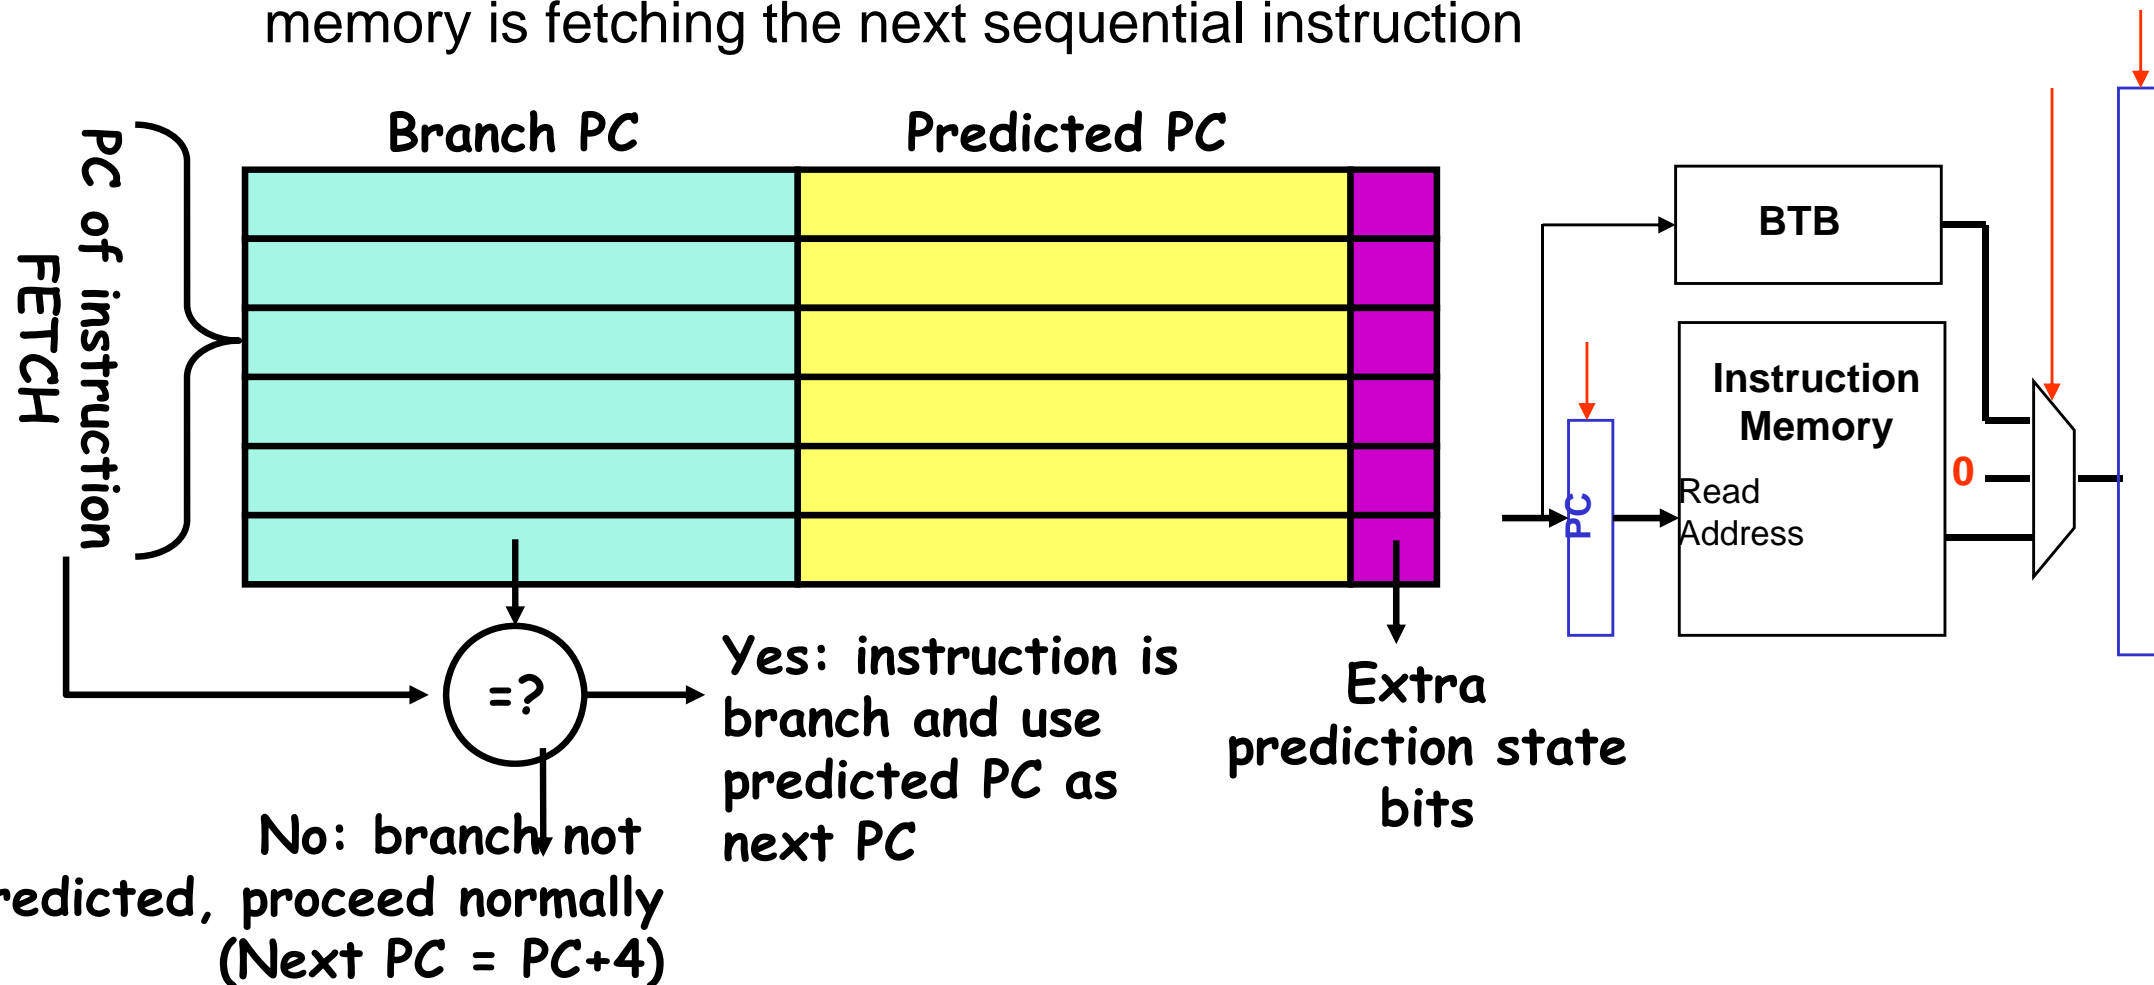
Loop: 1st loop instr
2nd loop instr
.
.
.
last loop instr
bne \$1,\$2,Loop
fall out instr

- BHT also stores the initial FSM state

Branch Target Buffer

- ❑ The BHT predicts *when* a branch is taken, but does not tell *where* its taken to!

- BTB can cache the branch taken *instruction* while the instruction memory is fetching the next sequential instruction

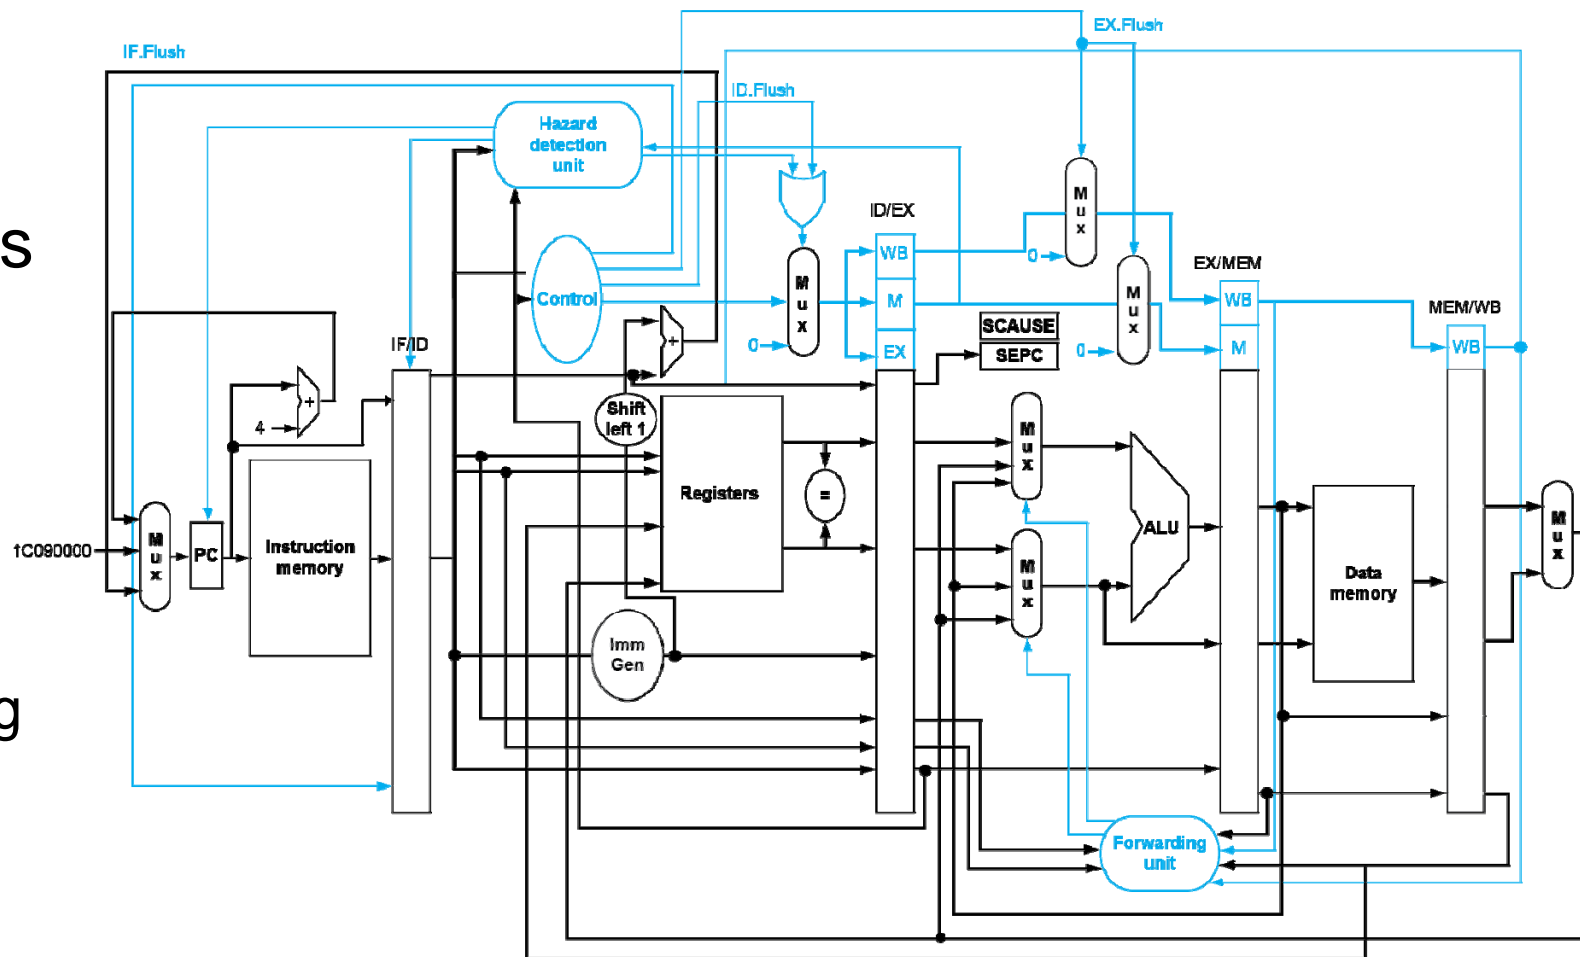


Exceptions on pipeline

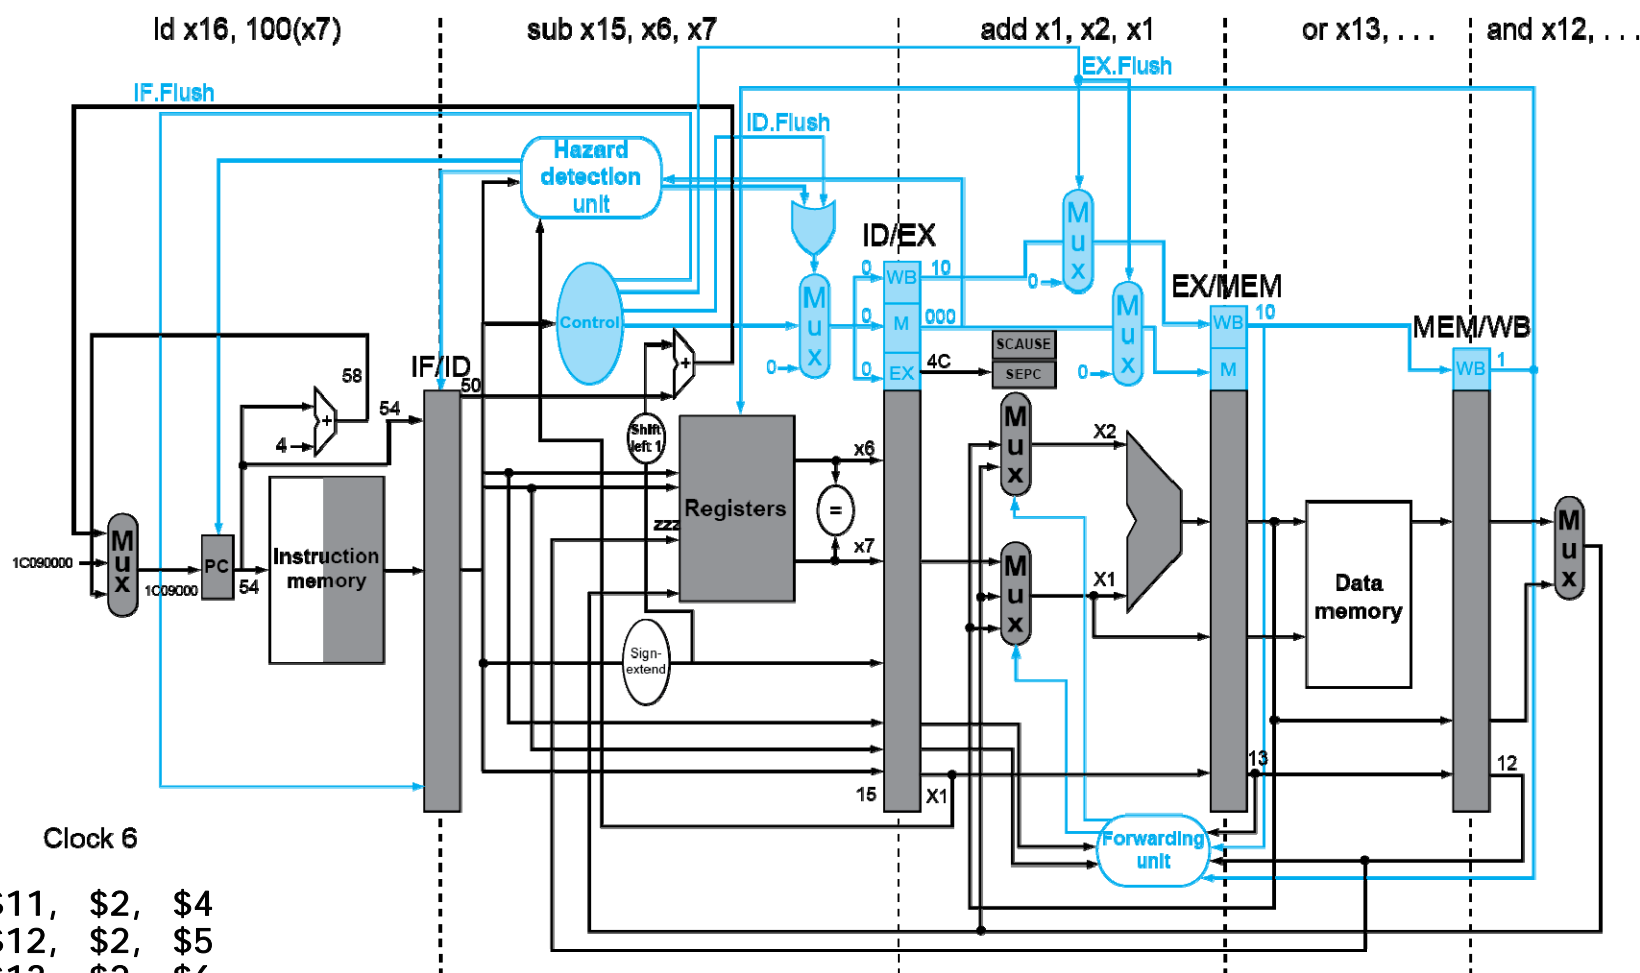
- ❑ Flush instructions that follow faulted instruction from pipeline
- ❑ Save address of offending instruction

- ❑ Precise interrupts

- complete prior instruction
- stop offending instruction in midstream
- flush all following instructions



Exception Example @ 6th cycle



Exception on add in

Clock 6

40	sub	\$11,	\$2,	\$4
44	and	\$12,	\$2,	\$5
48	or	\$13,	\$2,	\$6
4C	add	\$1,	\$2,	\$1
50	sl t	\$15,	\$6,	\$7
54	lw	\$16,	50(\$7)	

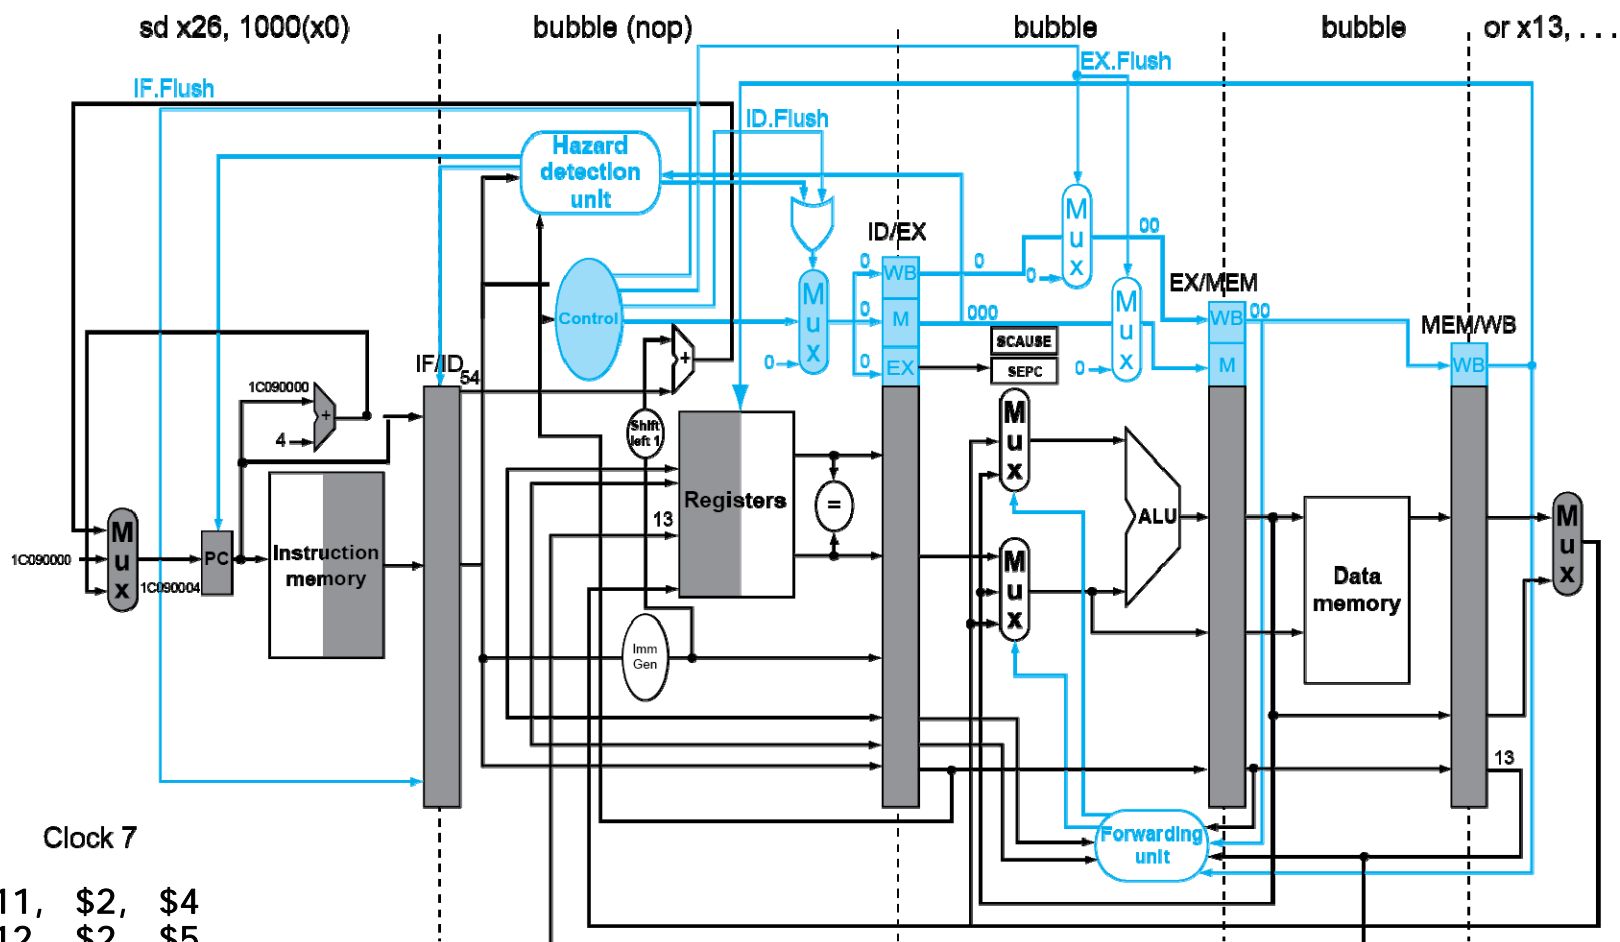
...

Handler

80000180	SW	\$25,	1000(\$0)
80000184	SW	\$26,	1004(\$0)

...

Exception Example @ 7th cycle



Exception on **add** in

Clock 7

40	sub	\$11,	\$2,	\$4
44	and	\$12,	\$2,	\$5
48	or	\$13,	\$2,	\$6
4C	add	\$1,	\$2,	\$1
50	sl t	\$15,	\$6,	\$7
54	lw	\$16,	50(\$7)	

...

Handler

80000180	sw	\$25,	1000(\$0)
80000184	sw	\$26,	1004(\$0)

...

Types of Parallelism

❑ **Instruction-level parallelism (ILP)** of a program – a measure of the average number of instructions in a program that a processor *might* be able to execute at the same time

- Mostly determined by the number of true (data) dependencies and procedural (control) dependencies in terms of instructions

❑ **Data-level parallelism (DLP)**

■ **To increase ILP**

❑ **Deeper pipeline**

Less work per stage shorter clock cycle

❑ **Multiple issue**

- Replicate pipeline stages multiple pipelines
- Start multiple instructions per clock cycle
- $CPI < 1$, so use Instructions Per Cycle (IPC)
- E.g., 4GHz 4-way multiple-issue
- 16 BIPS, peak $CPI = 0.25$, peak $IPC = 4$
- But dependencies reduce this in practice

```
DO   I = 1   TO   100  
    A[I] = A[I] + 1  
CONTINUE
```

Multiple-Issue Processor Styles

❑ Static multiple-issue processors (aka **VLIW**)

- Decisions on which instructions to execute simultaneously are being made statically (at compile time by the compiler)
- E.g., Intel Itanium and Itanium 2 for the IA-64 ISA – EPIC (Explicit Parallel Instruction Computer)
 - 128-bit “bundles” containing three instructions, each 41-bits plus a 5-bit template field (which specifies which FU each instruction needs)
 - Five functional units (IntALU, Mmedia, Dmem, FPALU, Branch)
 - Extensive support for speculation and predication

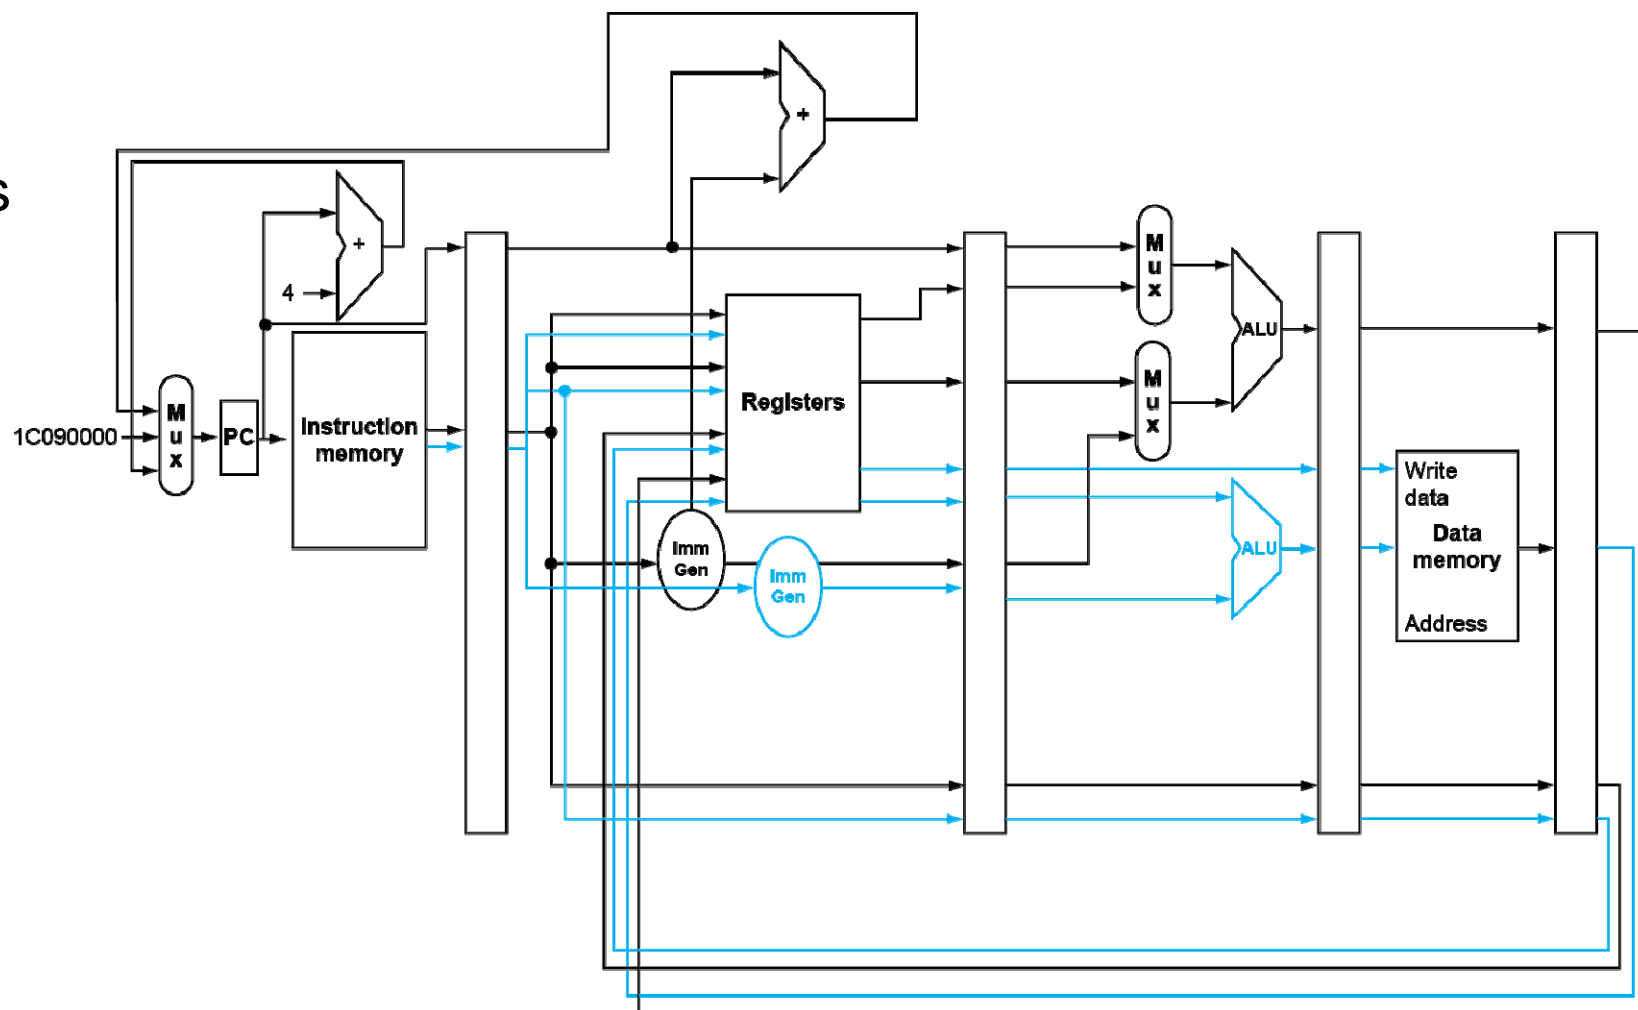
❑ Dynamic multiple-issue processors (aka **superscalar**)

- Decisions on which instructions to execute simultaneously (in the range of 2 to 8) are being made dynamically (at run time by the hardware)
- E.g., IBM Power series, Pentium 4, MIPS R10K, AMD Barcelona

A static two-issue datapath

□ What we add:

- two 32b instructions from Imem
- four read ports
- Two write ports
- Additional ALU



RISC-V with Static Dual Issue

- ❑ Two-issue packets
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

Scheduling Example

- ❑ Schedule this for dual-issue RISC-V

```
Loop: ld    x31, 0(x20)      // x31=array element
      add   x31, x31, x21    // add scalar in x21
      sd    x31, 0(x20)     // store result
      addi  x20, x20, -8     // decrement pointer
      blt   x22, x20, Loop   // branch if x22 < x20
```

	ALU/branch	Load/store	cycle
Loop:	nop	ld x31, 0(x20)	1
	addi x20, x20, -8	nop	2
	add x31, x31, x21	nop	3
	blt x22, x20, Loop	sd x31, 8(x20)	4


- $IPC = 5/4 = 1.25$ (c.f. peak $IPC = 2$)

Loop Unrolling

- ❑ Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
- ❑ Use different registers per replication
 - Called “register renaming”
 - Avoid loop-carried “anti-dependencies”
 - Store followed by a load of the same register
 - Aka “name dependence”
 - Reuse of a register name

```
DO  I = 1  TO  100
    A[I] = A[I] + 1
CONTINUE
```

```
DO  I = 1  TO  100 by 4
    A[I] = A[I] + 1
    A[I+1] = A[I+1] + 1
    A[I+2] = A[I+2] + 1
    A[I+3] = A[I+3] + 1
CONTINUE
```



Loop Unrolling Example

Loop: l d **x31**, 0(x20)
 add **x31**, **x31**, x21
 sd **x31**, 0(x20)
 addi **x20**, x20, -8
 bl t x22, **x20**, Loop

Loop: l d **x31**, 0(x20)
 add **x31**, **x31**, x21
 sd **x31**, 0(x20)
 addi **x20**, x20, -8
 bl t x22, **x20**, Loop

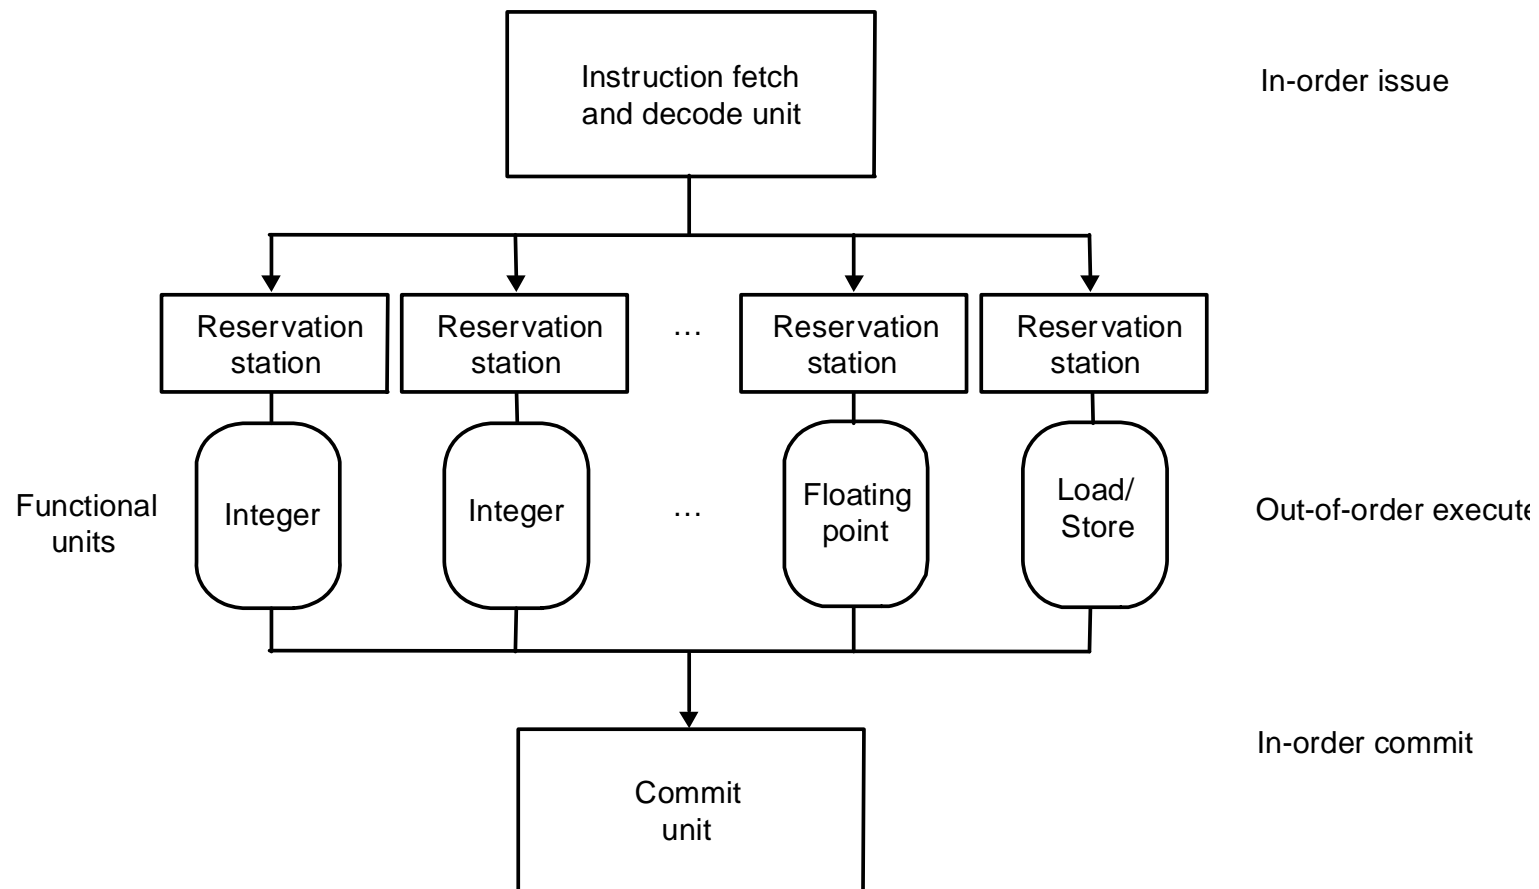
	ALU/branch	Load/store	cycle
Loop:	addi x20, x20, -32	l d x28 , 0(x20)	1
	nop	l d x29 , 24 (x20)	2
	add x28 , x28 , x21	l d x30 , 16 (x20)	3
	add x29 , x29 , x21	l d x31 , 8 (x20)	4
	add x30 , x30 , x21	sd x28 , 32 (x20)	5
	add x31 , x31 , x21	sd x29 , 24 (x20)	6
	nop	sd x30 , 16 (x20)	7
	bl t x22, x20, Loop	sd x31 , 8 (x20)	8

❑ $IPC = 14/8 = 1.75$

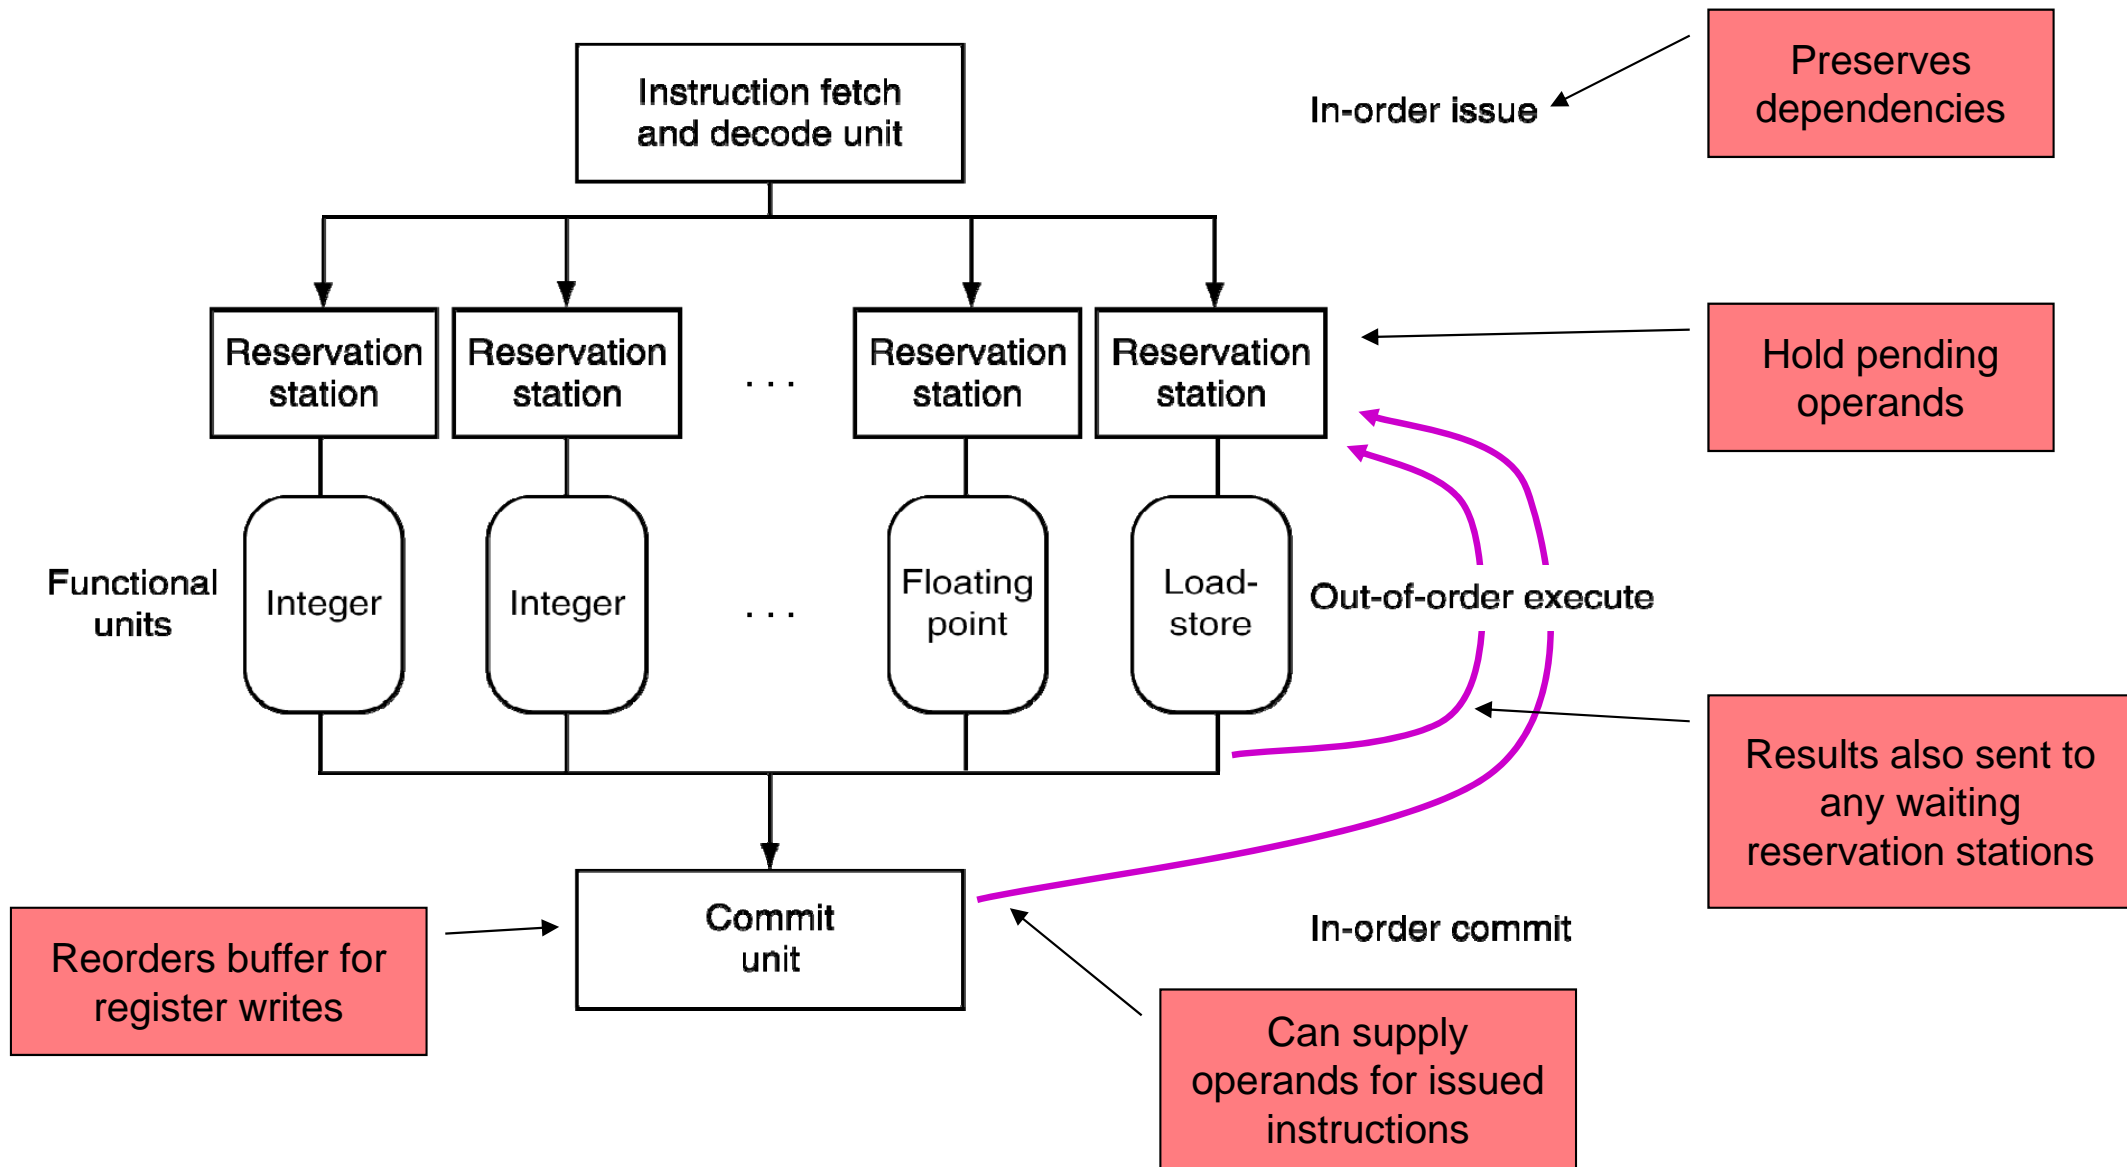
– Closer to 2, but at cost of registers and code size

Dynamic Pipeline Scheduling

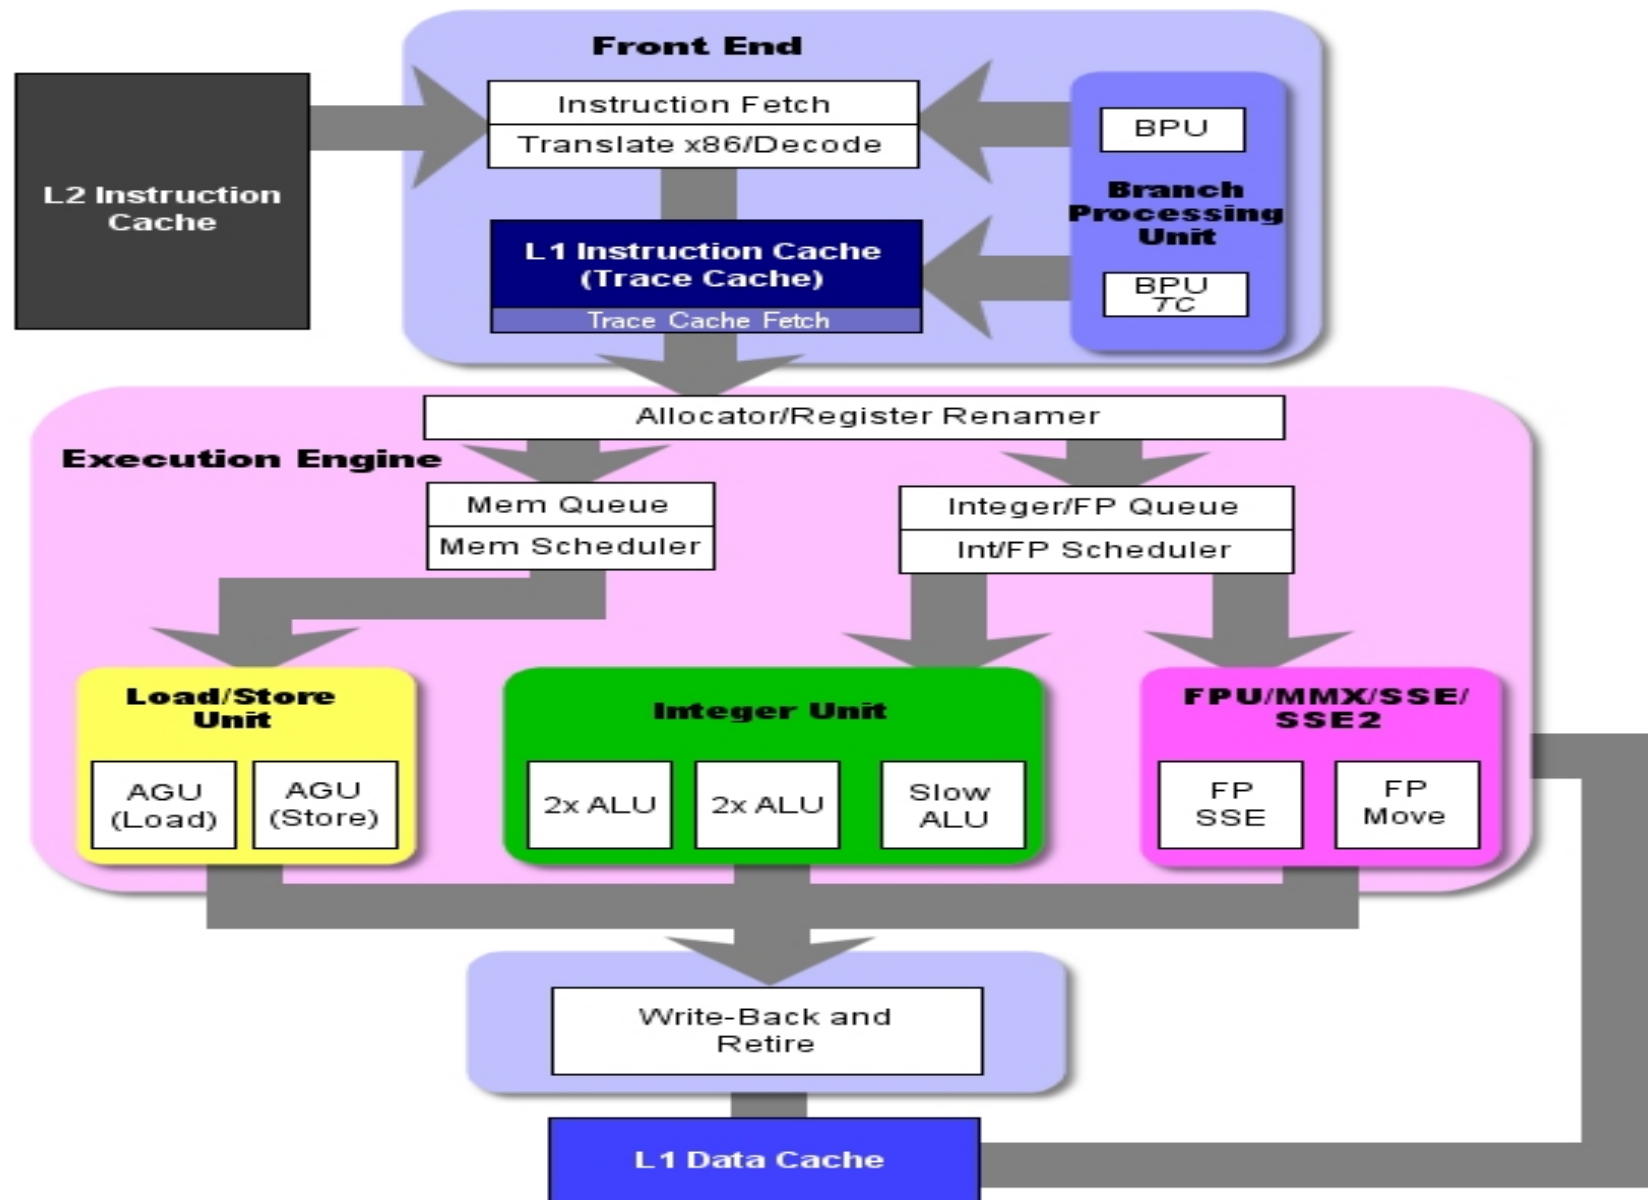
- ❑ hardware tries to find instructions to execute
- ❑ out of order execution is possible
- ❑ speculative execution and dynamic branch prediction



Superscalar by dynamically scheduling



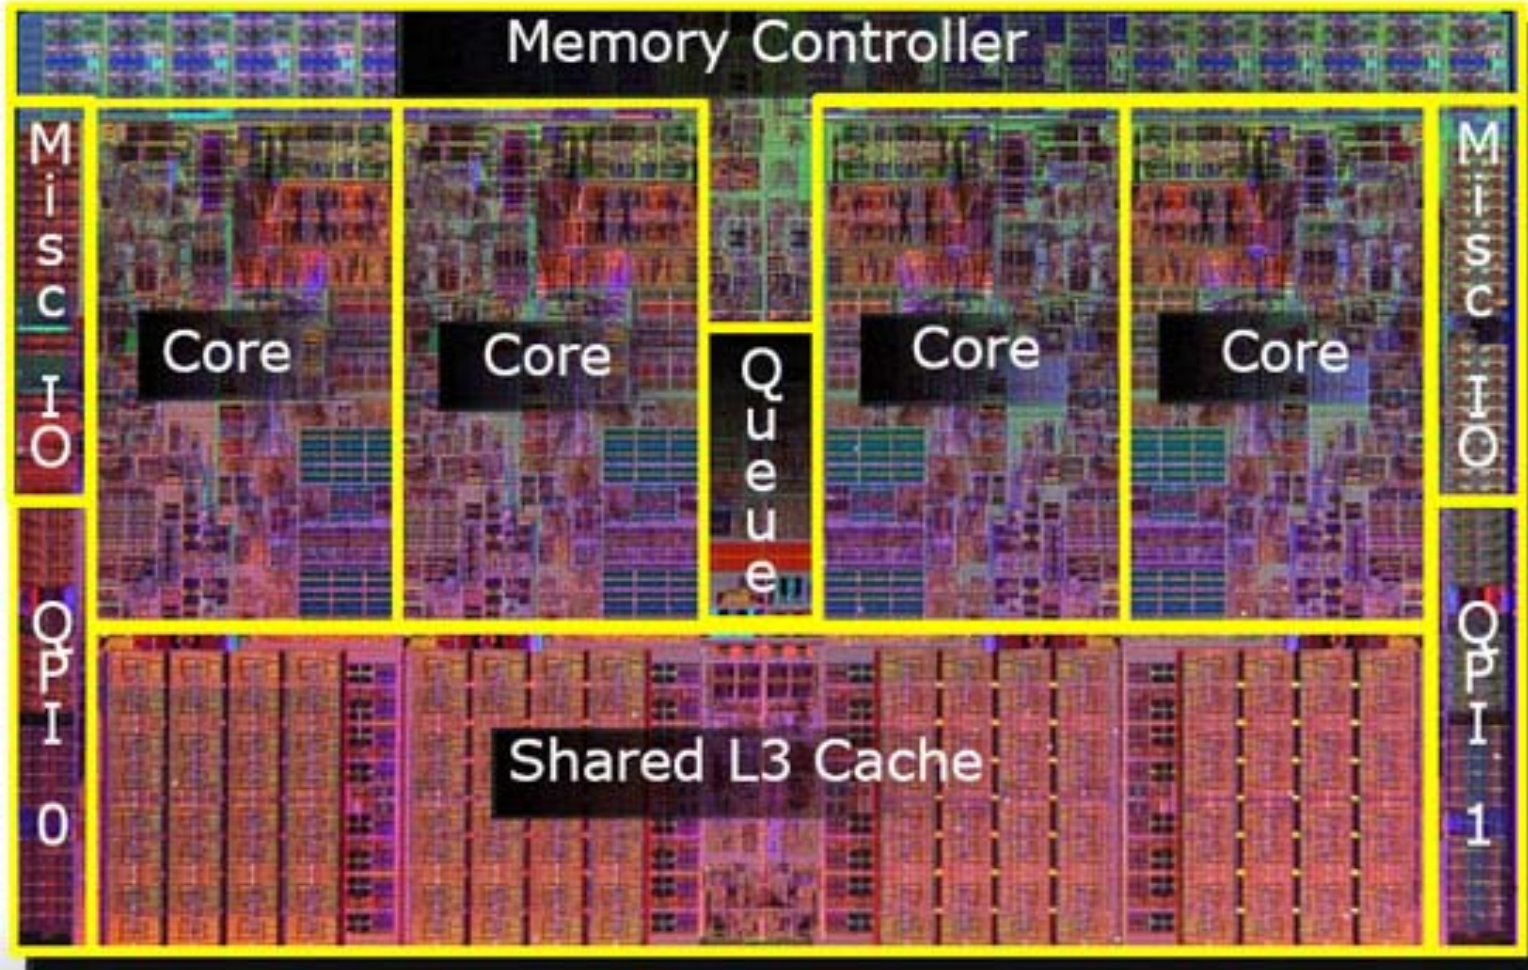
Intel Pentium IV



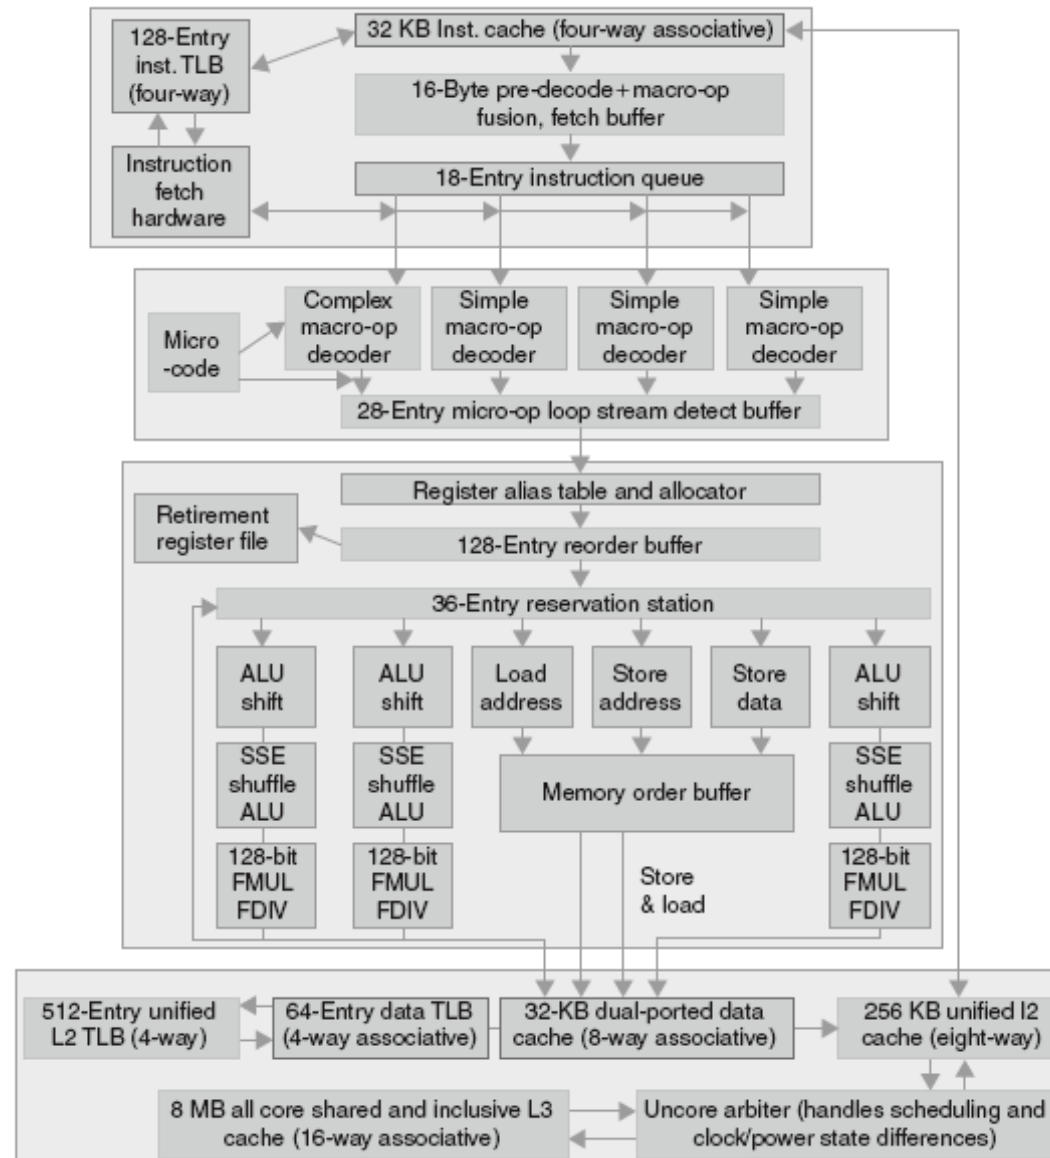
How about the most recent Intel i7?

❑ The new Core i7 processor

- 1366 pins
- 263 mm²
- 731 Million transistors



Core i7 Pipeline



Power Efficiency

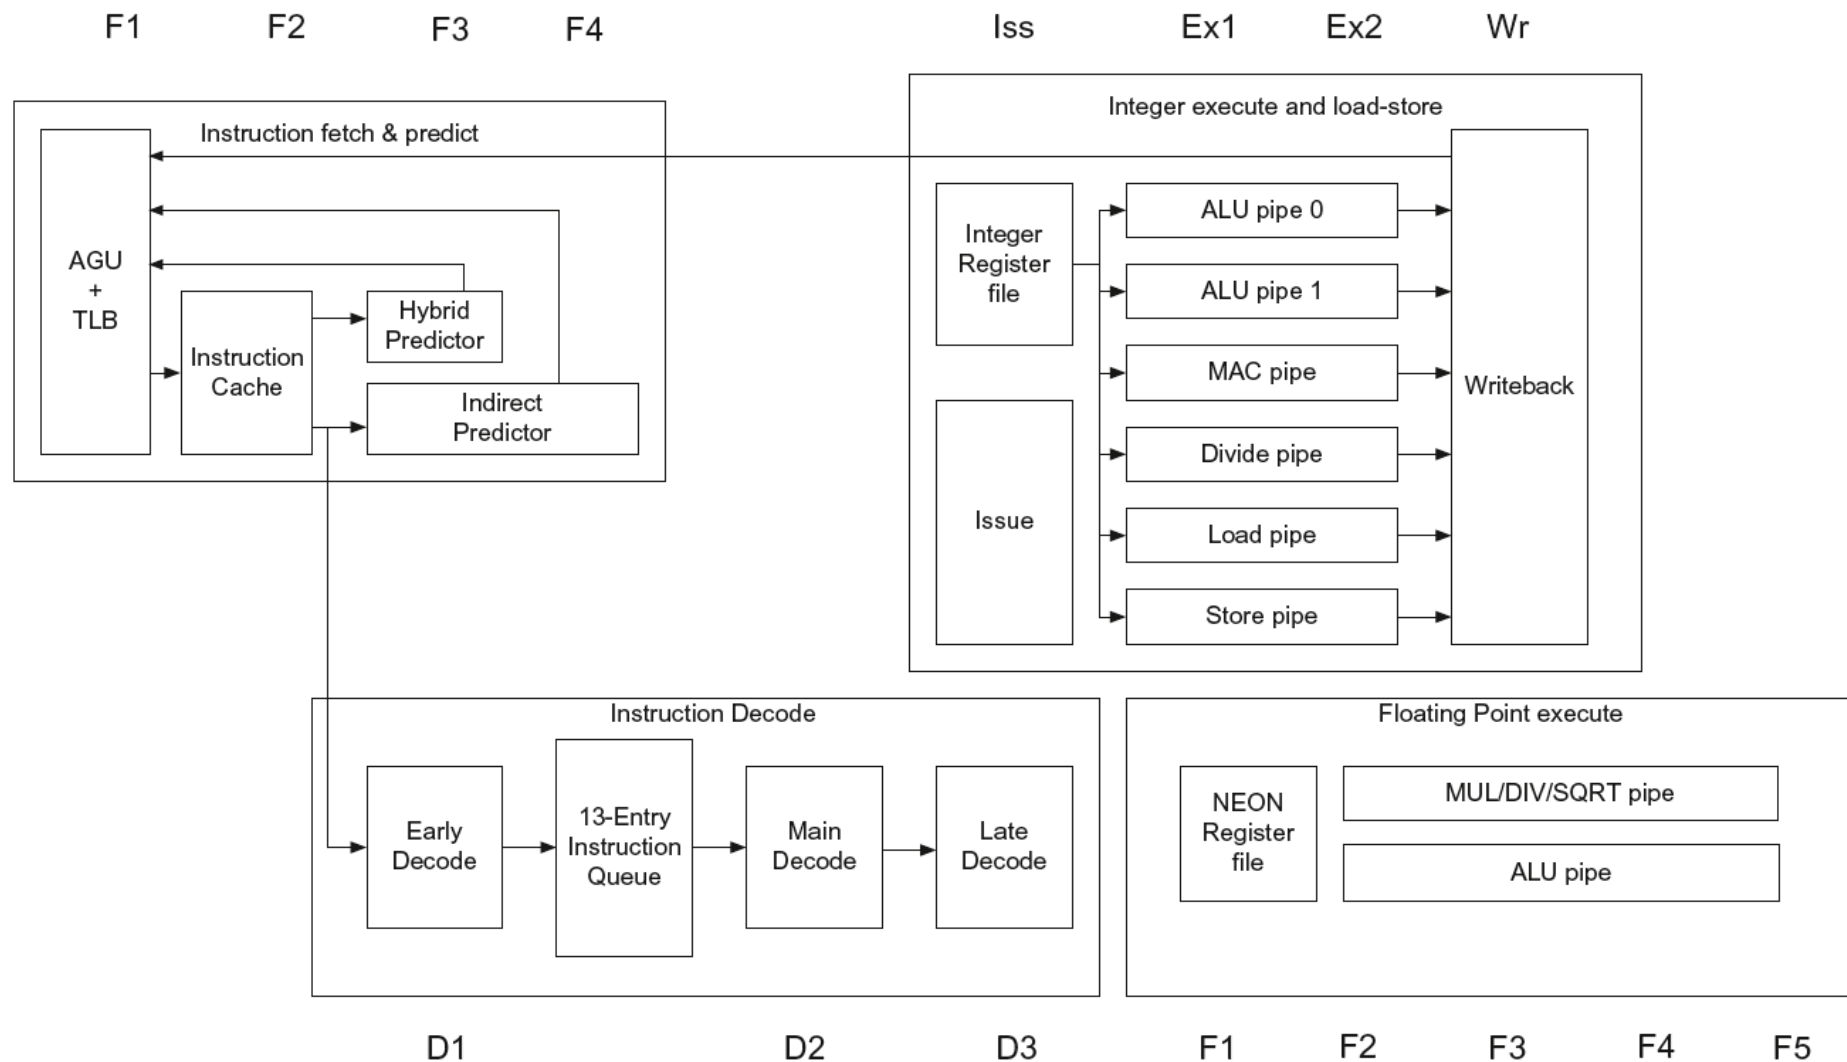
- ❑ Complexity of dynamic scheduling and speculations requires power
- ❑ Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

ARM Cortex A53 and Intel i7

Processor	ARM A53	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	100 milliWatts (1 core @ 1 GHz)	130 Watts
Clock rate	1.5 GHz	2.66 GHz
Cores/Chip	4 (configurable)	4
Floating point?	Yes	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	8	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	Hybrid	2-level
1 st level caches/core	16-64 KiB I, 16-64 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-2048 KiB	256 KiB (per core)
3 rd level caches (shared)	(platform dependent)	2-8 MB

ARM Cortex-A53 Pipeline



Fallacies

- ❑ Pipelining is easy (!)
 - The basic idea is easy
 - The devil is in the details
 - e.g., detecting data hazards
- ❑ Pipelining is independent of technology
 - So why haven't we always done pipelining?
 - More transistors make more advanced techniques feasible
 - Pipeline-related ISA design needs to take account of technology trends
 - e.g., predicated instructions

Pitfalls

- ❑ Poor ISA design can make pipelining harder
 - e.g., complex instruction sets (VAX, IA-32)
 - Significant overhead to make pipelining work
 - IA-32 micro-op approach
 - e.g., complex addressing modes
 - Register update side effects, memory indirection
 - e.g., delayed branches
 - Advanced pipelines have long delay slots